

In []: # Breast Cancer Wisconsin (Original) - Exploratory Data Analysis

This notebook performs exploratory data analysis (EDA) on the cleaned Breast Cancer Wisconsin (Original) dataset. The main goals are:

- To understand the distribution of each cytological feature.
- To analyze how each feature relates to the malignancy label.
- To visualize the results using correlation plots **and** boxplots.
- To train a simple Random Forest classifier **and** inspect feature importances.

In [3]: # Import basic packages

```
import os
import numpy as np
import pandas as pd
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```

Make plots look nicer

```
plt.rcParams["figure.figsize"] = (8, 5)
plt.rcParams["axes.titlesize"] = 14
plt.rcParams["axes.labelsize"] = 12
```

Set path to the cleaned CSV file

```
DATA_PATH = "/Users/shijietang/Desktop/143project/Analysis/breast_cancer_wiscons
```

```
DATA_PATH
```

Out[3]: '/Users/shijietang/Desktop/143project/Analysis/breast_cancer_wisconsin_original_cleaned.csv'

In [4]: # Load the dataset

```
df = pd.read_csv(DATA_PATH)
```

```
print(f"Data shape: {df.shape[0]} rows, {df.shape[1]} columns")
df.head()
```

Data shape: 675 rows, 11 columns

Out[4]:

| | Clump Thickness | Uniformity of Cell Size | Uniformity of Cell Shape | Marginal Adhesion | Single Epithelial Cell Size | Bare Nuclei | Bland Chromatin | Nu |
|---|-----------------|-------------------------|--------------------------|-------------------|-----------------------------|-------------|-----------------|----|
| 0 | 5 | 1 | 1 | 1 | 2 | 1.0 | 3 | |
| 1 | 5 | 4 | 4 | 5 | 7 | 10.0 | 3 | |
| 2 | 3 | 1 | 1 | 1 | 2 | 2.0 | 3 | |
| 3 | 6 | 8 | 8 | 1 | 3 | 4.0 | 3 | |
| 4 | 4 | 1 | 1 | 3 | 2 | 1.0 | 3 | |

In [5]: print("Columns:", df.columns.tolist())
print("\nData types:\n", df.dtypes)

```
print("\nMissing values per column:\n")
print(df.isna().sum())
```

```
df.describe()
```

Columns: ['Clump Thickness', 'Uniformity of Cell Size', 'Uniformity of Cell Shape', 'Marginal Adhesion', 'Single Epithelial Cell Size', 'Bare Nuclei', 'Bland Chromatin', 'Normal Nucleoli', 'Mitoses', 'Diagnosis', 'Is_Malignant']

Data types:

| | |
|-----------------------------|---------|
| Clump Thickness | int64 |
| Uniformity of Cell Size | int64 |
| Uniformity of Cell Shape | int64 |
| Marginal Adhesion | int64 |
| Single Epithelial Cell Size | int64 |
| Bare Nuclei | float64 |
| Bland Chromatin | int64 |
| Normal Nucleoli | int64 |
| Mitoses | int64 |
| Diagnosis | int64 |
| Is_Malignant | int64 |

dtype: object

Missing values per column:

| | |
|-----------------------------|---|
| Clump Thickness | 0 |
| Uniformity of Cell Size | 0 |
| Uniformity of Cell Shape | 0 |
| Marginal Adhesion | 0 |
| Single Epithelial Cell Size | 0 |
| Bare Nuclei | 0 |
| Bland Chromatin | 0 |
| Normal Nucleoli | 0 |
| Mitoses | 0 |
| Diagnosis | 0 |
| Is_Malignant | 0 |

dtype: int64

Out[5]:

| | Clump Thickness | Uniformity of Cell Size | Uniformity of Cell Shape | Marginal Adhesion | Single Epithelial Cell Size | Bare Nuclei | C |
|--------------|-----------------|-------------------------|--------------------------|-------------------|-----------------------------|-------------|------------|
| count | 675.000000 | 675.000000 | 675.000000 | 675.000000 | 675.000000 | 675.000000 | 675.000000 |
| mean | 4.451852 | 3.146667 | 3.208889 | 2.848889 | 3.229630 | 3.537778 | 3.537778 |
| std | 2.820859 | 3.055005 | 2.976552 | 2.875917 | 2.208497 | 3.637871 | 3.637871 |
| min | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 |
| 25% | 2.000000 | 1.000000 | 1.000000 | 1.000000 | 2.000000 | 1.000000 | 1.000000 |
| 50% | 4.000000 | 1.000000 | 1.000000 | 1.000000 | 2.000000 | 1.000000 | 1.000000 |
| 75% | 6.000000 | 5.000000 | 5.000000 | 4.000000 | 4.000000 | 6.000000 | 6.000000 |
| max | 10.000000 | 10.000000 | 10.000000 | 10.000000 | 10.000000 | 10.000000 | 10.000000 |



In [6]:

```
counts = df["Is_Malignant"].value_counts().sort_index()
ratios = df["Is_Malignant"].value_counts(normalize=True).sort_index()

print("Class distribution for Is_Malignant:")
for c in counts.index:
```

```
label = "Benign (0)" if c == 0 else "Malignant (1)"
print(f"{label}: {counts[c]} samples ({ratios[c] * 100:.2f}%)")
```

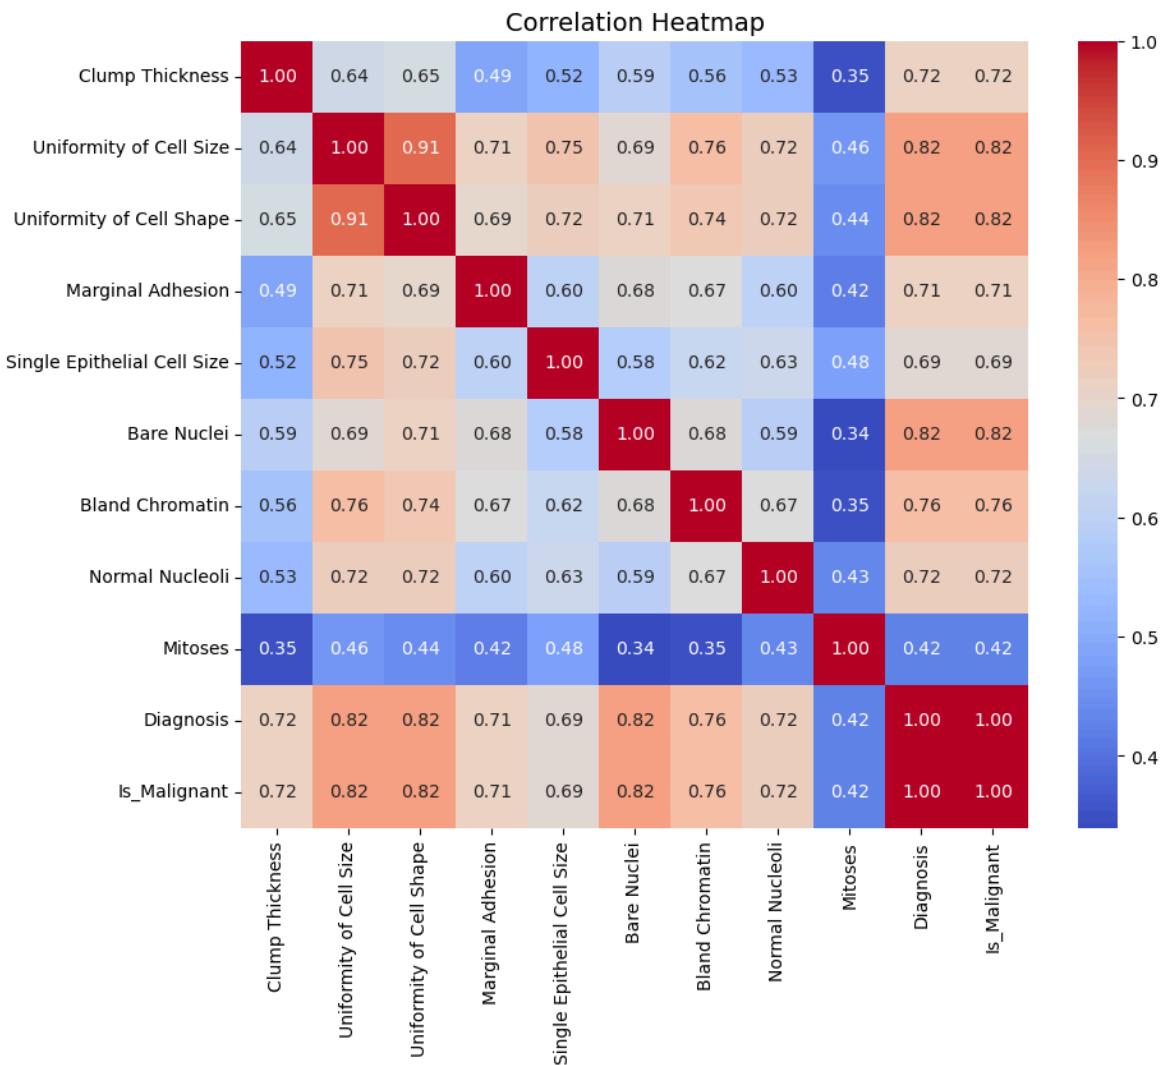
Class distribution for Is_Malignant:

Benign (0): 439 samples (65.04%)

Malignant (1): 236 samples (34.96%)

```
In [7]: numeric_cols = df.select_dtypes(include=[np.number]).columns.tolist()
corr_matrix = df[numeric_cols].corr()

plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, fmt=".2f", cmap="coolwarm", square=True)
plt.title("Correlation Heatmap")
plt.show()
```

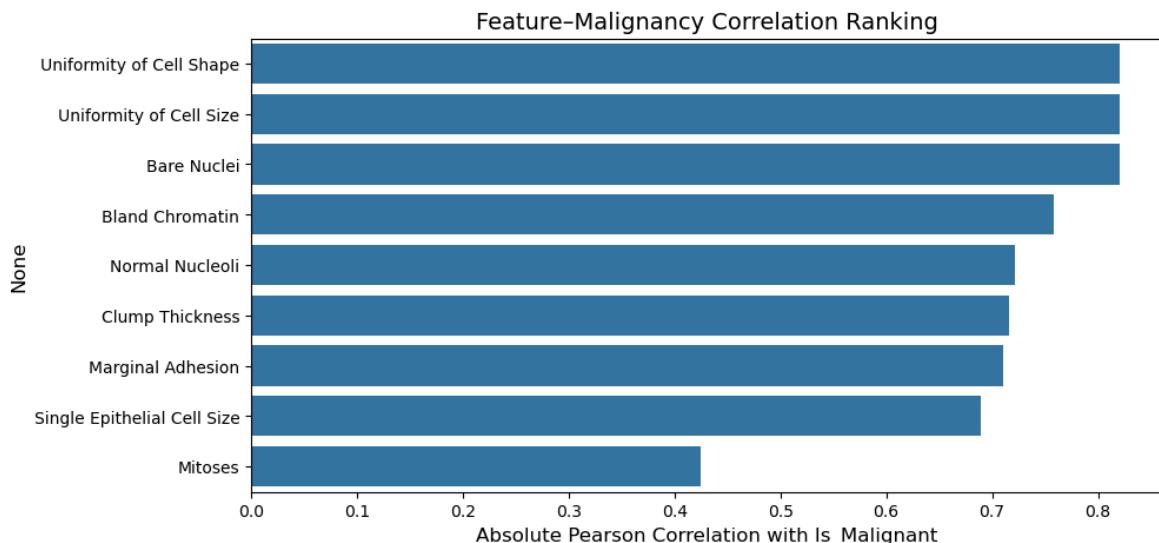


```
In [8]: corr_with_target = df[numeric_cols].corr()["Is_Malignant"].sort_values(ascending=False)
```

```
Out[8]: Is_Malignant      1.000000
Diagnosis          1.000000
Uniformity of Cell Shape 0.820688
Uniformity of Cell Size   0.820663
Bare Nuclei        0.820533
Bland Chromatin     0.758252
Normal Nucleoli    0.722020
Clump Thickness     0.715714
Marginal Adhesion    0.710985
Single Epithelial Cell Size 0.689544
Mitoses            0.424399
Name: Is_Malignant, dtype: float64
```

```
In [9]: # Remove the target & original diagnosis code from feature ranking
corr_features_only = corr_with_target.drop(["Is_Malignant", "Diagnosis"])

plt.figure(figsize=(10, 5))
sns.barplot(x=corr_features_only.abs(), y=corr_features_only.index)
plt.xlabel("Absolute Pearson Correlation with Is_Malignant")
plt.title("Feature-Malignancy Correlation Ranking")
plt.show()
```



```
In [10]: df_plot = df.copy()
df_plot["Diagnosis_Label"] = df_plot["Is_Malignant"].map({0: "Benign (0)", 1: "Malignant (1)"})

features = [col for col in numeric_cols if col not in ["Is_Malignant", "Diagnosis"]]
features
```

```
Out[10]: ['Clump Thickness',
          'Uniformity of Cell Size',
          'Uniformity of Cell Shape',
          'Marginal Adhesion',
          'Single Epithelial Cell Size',
          'Bare Nuclei',
          'Bland Chromatin',
          'Normal Nucleoli',
          'Mitoses']
```

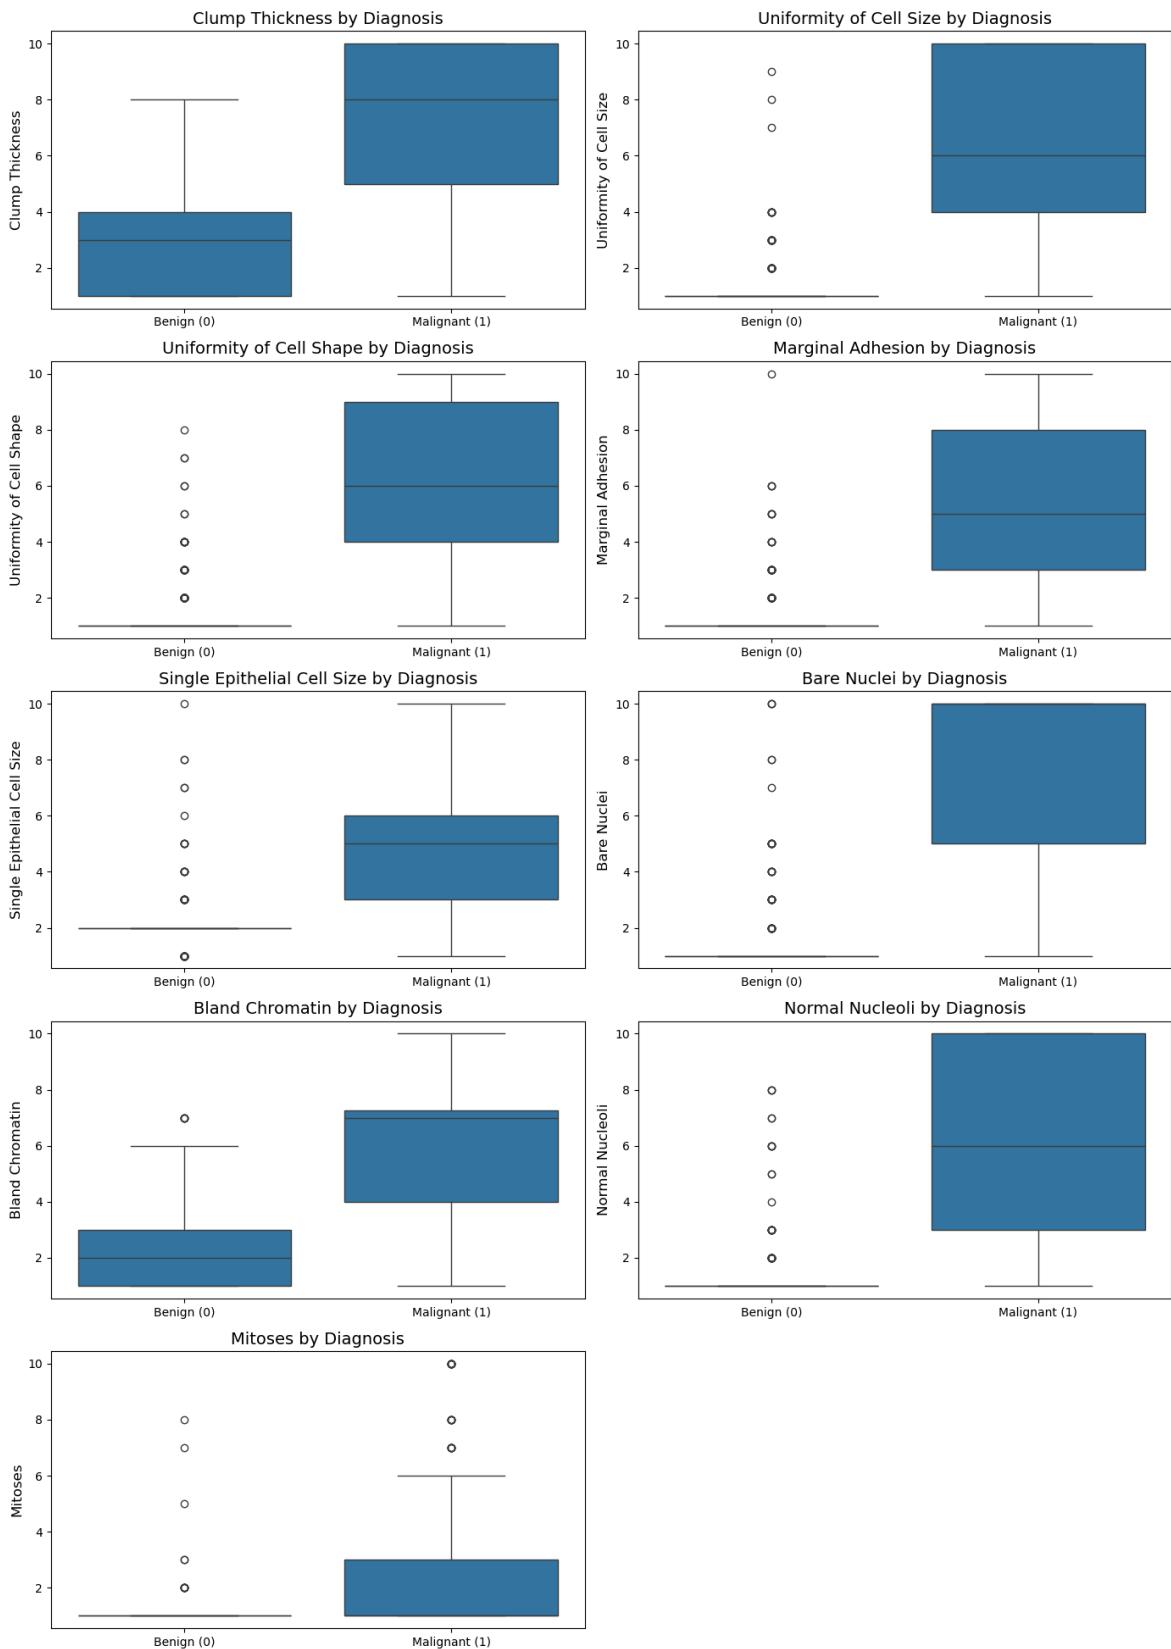
```
In [11]: # Multi-panel boxplots
n_features = len(features)
n_cols = 2
n_rows = int(np.ceil(n_features / n_cols))
```

```
plt.figure(figsize=(14, 4 * n_rows))

for i, feature in enumerate(features, start=1):
    plt.subplot(n_rows, n_cols, i)
    sns.boxplot(data=df_plot, x="Diagnosis_Label", y=feature)
    plt.title(f"{feature} by Diagnosis")
    plt.xlabel("")
    plt.ylabel(feature)

plt.suptitle("Feature Distributions by Diagnosis", fontsize=16, y=1.02)
plt.tight_layout()
plt.show()
```

Feature Distributions by Diagnosis



```
In [12]: means = df.groupby("Is_Malignant")[features].mean()
means
```

Out[12]:

| | Clump Thickness | Uniformity of Cell Size | Uniformity of Cell Shape | Marginal Adhesion | Single Epithelial Cell Size | Bare Nuclei | Cl |
|---------------------|-----------------|-------------------------|--------------------------|-------------------|-----------------------------|-------------|----|
| Is_Malignant | | | | | | | |
| 0 | 2.972665 | 1.309795 | 1.419134 | 1.350797 | 2.113895 | 1.350797 | |
| 1 | 7.203390 | 6.563559 | 6.538136 | 5.635593 | 5.305085 | 7.605932 | |

In [13]:

```
(means.loc[1] - means.loc[0]).sort_values(ascending=False)
```

Out[13]:

| | |
|-----------------------------|----------|
| Bare Nuclei | 6.255135 |
| Uniformity of Cell Size | 5.253764 |
| Uniformity of Cell Shape | 5.119001 |
| Normal Nucleoli | 4.638305 |
| Marginal Adhesion | 4.284796 |
| Clump Thickness | 4.230725 |
| Bland Chromatin | 3.899087 |
| Single Epithelial Cell Size | 3.191190 |
| Mitoses | 1.548348 |

dtype: float64

In [14]:

```
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import roc_auc_score, classification_report

X = df[features].values
y = df["Is_Malignant"].values

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, stratify=y, random_state=42
)

rf = RandomForestClassifier(
    n_estimators=200,
    class_weight="balanced",
    random_state=42,
)

rf.fit(X_train, y_train)

y_prob = rf.predict_proba(X_test)[:, 1]
y_pred = rf.predict(X_test)

print("ROC-AUC:", roc_auc_score(y_test, y_prob))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

ROC-AUC: 0.994771660264618

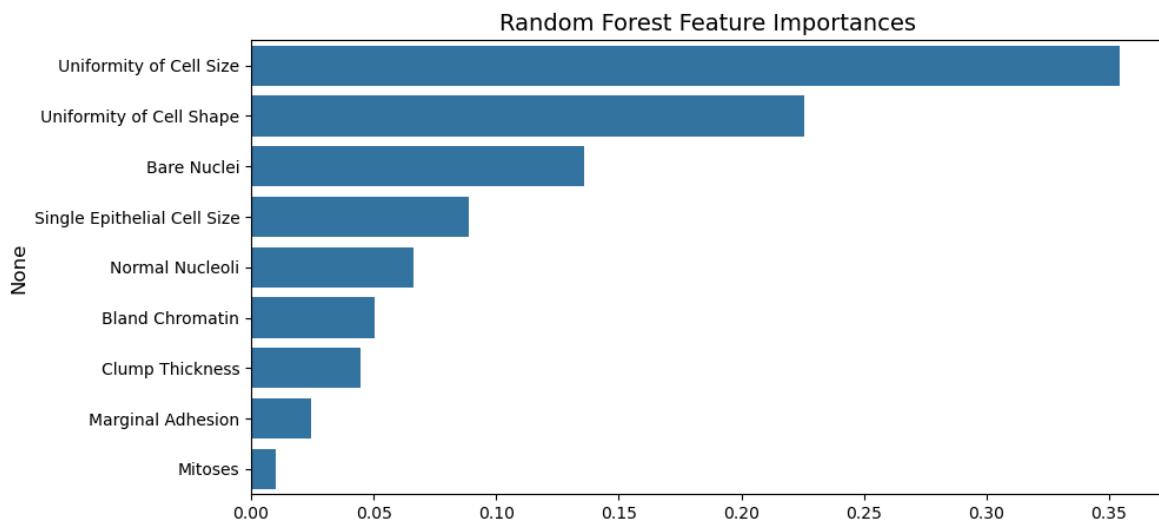
Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.96 | 0.98 | 0.97 | 132 |
| 1 | 0.96 | 0.93 | 0.94 | 71 |
| accuracy | | | 0.96 | 203 |
| macro avg | 0.96 | 0.95 | 0.96 | 203 |
| weighted avg | 0.96 | 0.96 | 0.96 | 203 |

```
In [15]: importances = pd.Series(rf.feature_importances_, index=features).sort_values(ascending=True)
```

```
Out[15]: Uniformity of Cell Size      0.354379
Uniformity of Cell Shape      0.225554
Bare Nuclei                  0.135997
Single Epithelial Cell Size  0.088600
Normal Nucleoli              0.066054
Bland Chromatin              0.050582
Clump Thickness               0.044595
Marginal Adhesion             0.024356
Mitoses                      0.009884
dtype: float64
```

```
In [16]: plt.figure(figsize=(10, 5))
sns.barplot(x=importances.values, y=importances.index)
plt.title("Random Forest Feature Importances")
plt.show()
```



```
In [ ]:
```