

Exploratory Data Analysis – Breast Cancer Coimbra Dataset

This notebook performs exploratory data analysis (EDA) on the **Breast Cancer Coimbra** dataset.

- Goal: Explore clinical / metabolic features and their relationship with breast cancer status.
- Target variable: `Classification`
 - 0 = Healthy / Control
 - 1 = Patient / Malignant

This analysis includes:

1. Data loading and basic checks
2. Univariate analysis (distributions & summary statistics)
3. Bivariate analysis (feature vs target)
4. Correlation analysis
5. Outlier inspection
6. Scaling exploration
7. Simple feature-importance measures

More details and interpretations will be noted on later sections at where the corresponding data was been extracted. A final summary is noted at the end as well.

```
In [1]: #import Lib
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
from sklearn.feature_selection import mutual_info_classif, f_classif
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
```

1. Dataset Overview

Shape & Columns

- **Rows:** 4,000
- **Feature columns:** Age, BMI, Glucose, Insulin, HOMA, Leptin, Adiponectin, Resistin, MCP.1
- **Target:** `Classification` (0 or 1)

Data Types

All 9 features are numerical (`int64` or `float64`) and fully non-null.
No missing values were detected in any column.

Class Distribution

Class	Count	Proportion
0 – Healthy	1784	44.6%
1 – Patient	2216	55.4%

The dataset shows a **slight class imbalance**, with more patient samples than healthy ones.

A count plot confirms this visually.

```
In [5]: #Load data
# Update the path if needed
df = pd.read_csv(r'D:\learning UCSD\ece143\Group_Project\breast_cancer_coimbra_c
print("Shape:", df.shape)

Shape: (4000, 10)

In [4]: print("DataFrame Info:")
print(df.info())
print("\nSummary statistics:")
display(df.describe().T)

print("\nMissing values per column:")
print(df.isna().sum())

DataFrame Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4000 entries, 0 to 3999
Data columns (total 10 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   Age               4000 non-null    int64  
 1   BMI               4000 non-null    float64 
 2   Glucose           4000 non-null    int64  
 3   Insulin           4000 non-null    float64 
 4   HOMA              4000 non-null    float64 
 5   Leptin             4000 non-null    float64 
 6   Adiponectin       4000 non-null    float64 
 7   Resistin          4000 non-null    float64 
 8   MCP.1             4000 non-null    float64 
 9   Classification    4000 non-null    int64  
dtypes: float64(7), int64(3)
memory usage: 312.6 KB
None

Summary statistics:
```

	count	mean	std	min	25%	50%
Age	4000.0	56.210750	17.809650	32.000000	39.000000	56.000000
BMI	4000.0	27.422280	4.413884	20.690751	23.079053	27.558485
Glucose	4000.0	113.876500	25.837795	76.000000	76.000000	131.000000
Insulin	4000.0	8.654001	6.435160	2.821000	4.421750	5.818000
HOMA	4000.0	2.024332	1.625638	0.590033	0.970090	1.373842
Leptin	4000.0	25.137737	15.096446	6.831900	12.712750	19.805050
Adiponectin	4000.0	9.364896	4.674244	3.192272	5.580210	8.286938
Resistin	4000.0	13.053667	7.454424	4.190320	7.022095	10.692780
MCP.1	4000.0	512.183456	253.279015	137.488000	299.665000	482.308000
Classification	4000.0	0.554000	0.497138	0.000000	0.000000	1.000000

◀ ▶

Missing values per column:

```
Age          0
BMI          0
Glucose      0
Insulin      0
HOMA          0
Leptin        0
Adiponectin  0
Resistin      0
MCP.1         0
Classification 0
dtype: int64
```

```
In [10]: TARGET_COL = "Classification"
feature_cols = [col for col in df.columns if col != TARGET_COL]

print("Target column:", TARGET_COL)
print("Feature columns:", feature_cols)

class_counts = df[TARGET_COL].value_counts().sort_index()
print("\nClass distribution:")
print(class_counts)

class_proportions = class_counts / len(df)
print("\nClass proportions:")
print(class_proportions)
```

```
Target column: Classification
Feature columns: ['Age', 'BMI', 'Glucose', 'Insulin', 'HOMA', 'Leptin', 'Adiponectin', 'Resistin', 'MCP.1']

Class distribution:
Classification
0    1784
1    2216
Name: count, dtype: int64

Class proportions:
Classification
0    0.446
1    0.554
Name: count, dtype: float64
```

2. Univariate Feature Distributions

Histograms and KDE curves were generated for each feature.

Key Observations

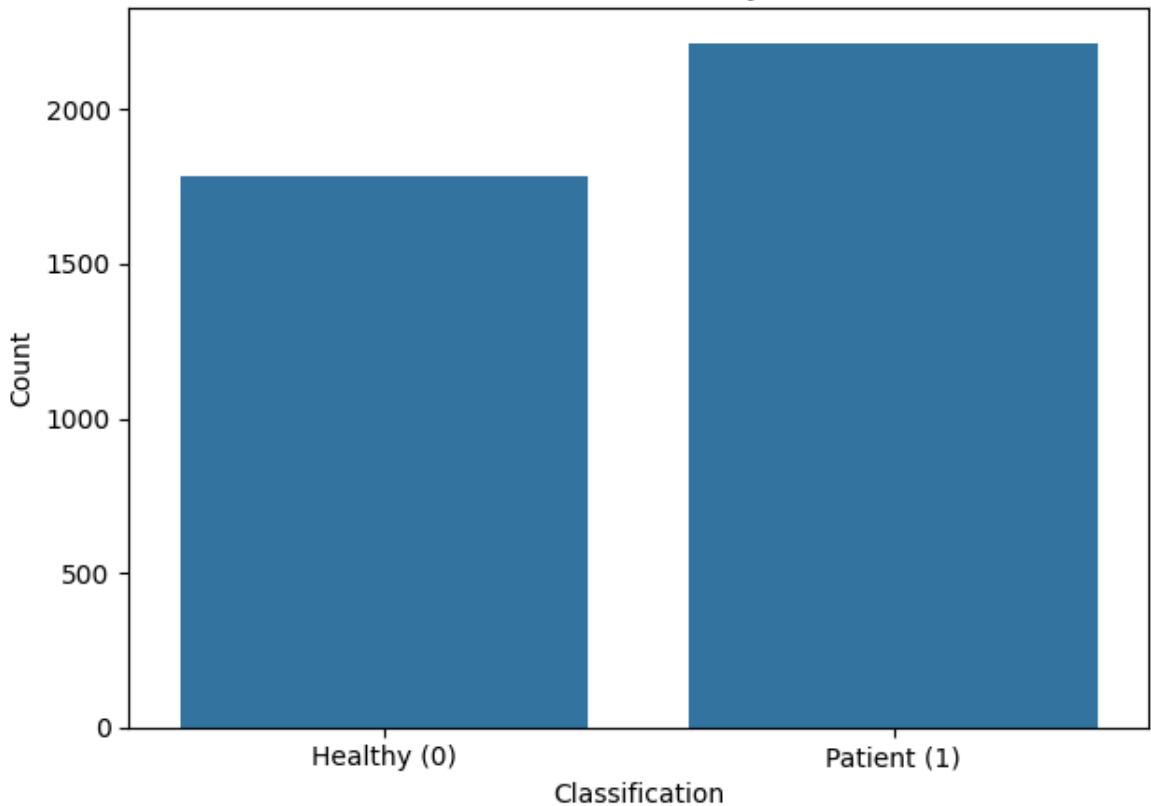
- **Age:** Broad distribution centered around ~56 years.
- **BMI:** Range ~21–36, mild multimodality.
- **Glucose:** Clearly **bimodal** with peaks around 76 and 131 mg/dL.
- **Insulin, HOMA, Resistin:** Strong **right skew** with long tails.
- **Leptin and Adiponectin:** Mild right skew.

Interpretation

Several metabolic markers show skewed distributions, suggesting that **log-transforms or robust scaling** may help stabilize variance in modeling.

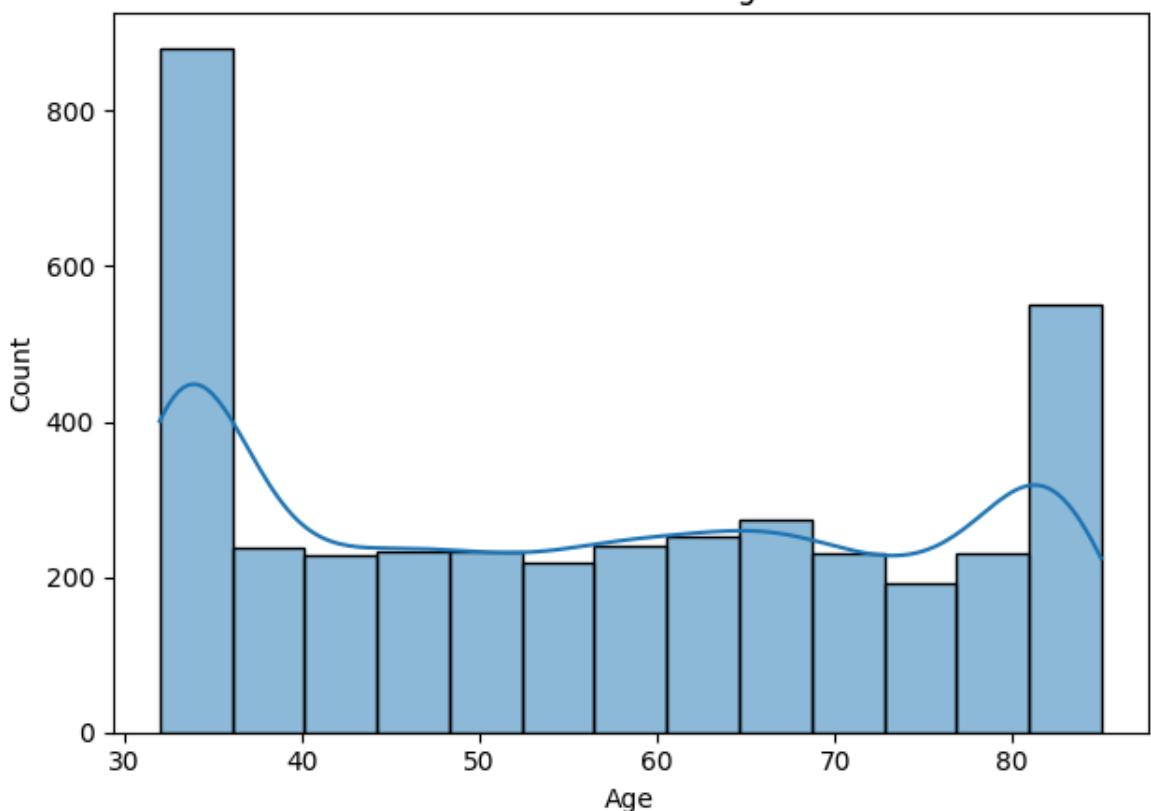
```
In [30]: fig, ax = plt.subplots()
sns.countplot(x=TARGET_COL, data=df, ax=ax)
ax.set_xticks([0, 1])
ax.set_xticklabels(["Healthy (0)", "Patient (1)"])
ax.set_title("Class Distribution: Healthy vs Patient")
ax.set_xlabel("Classification")
ax.set_ylabel("Count")
plt.tight_layout()
plt.show()
```

Class Distribution: Healthy vs Patient

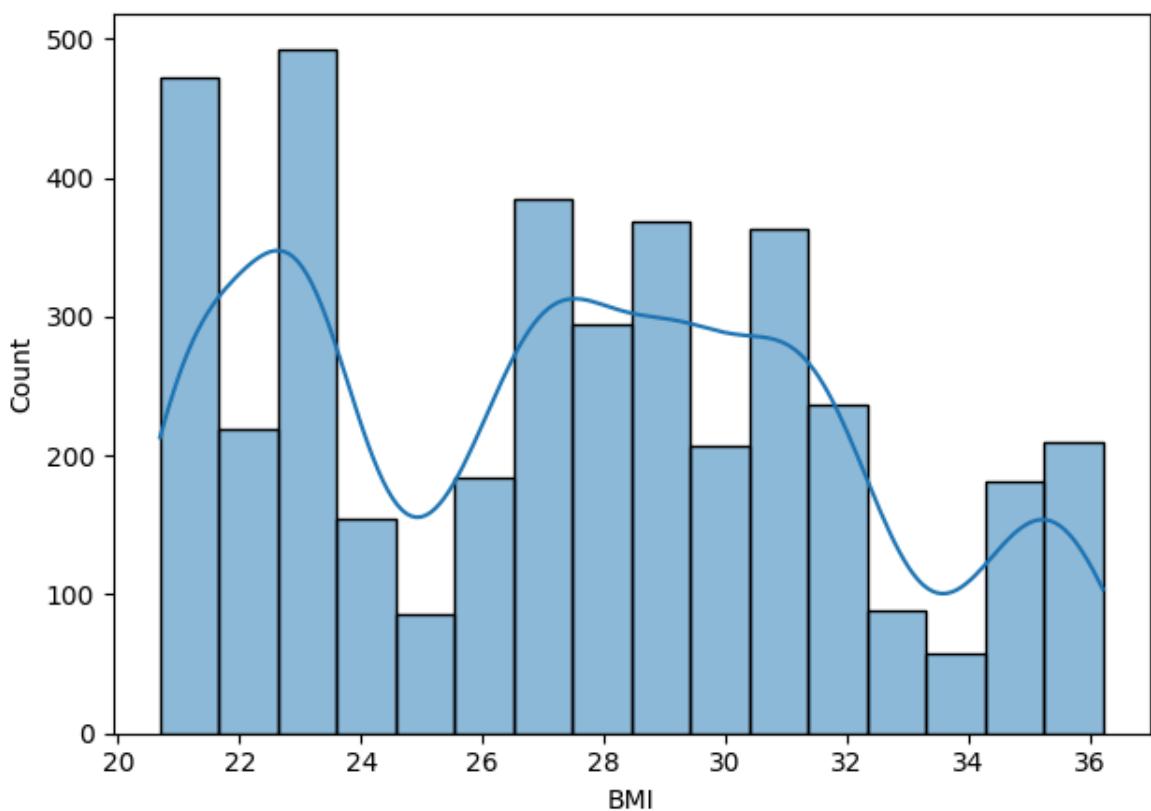


```
In [12]: # Histograms/KDE (overall)
for col in feature_cols:
    fig, ax = plt.subplots()
    sns.histplot(df[col], kde=True, ax=ax)
    ax.set_title(f"Distribution of {col}")
    plt.tight_layout()
    plt.show()
```

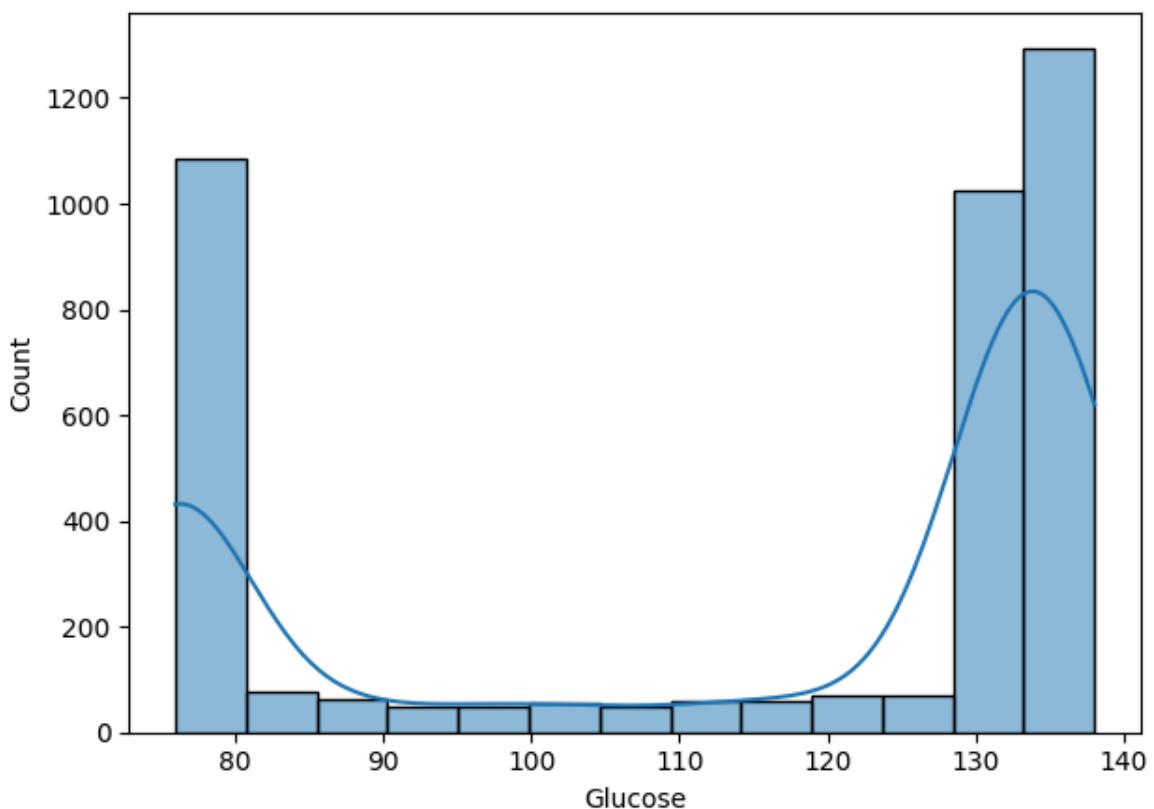
Distribution of Age



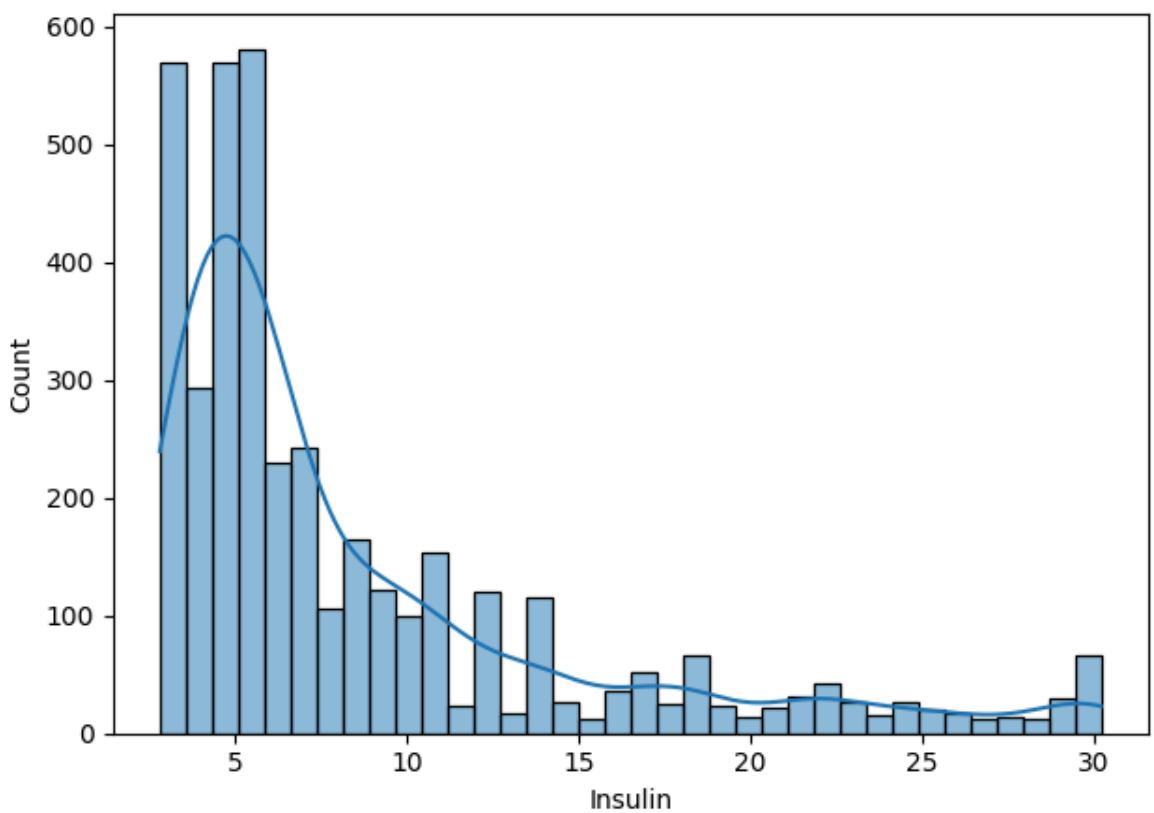
Distribution of BMI



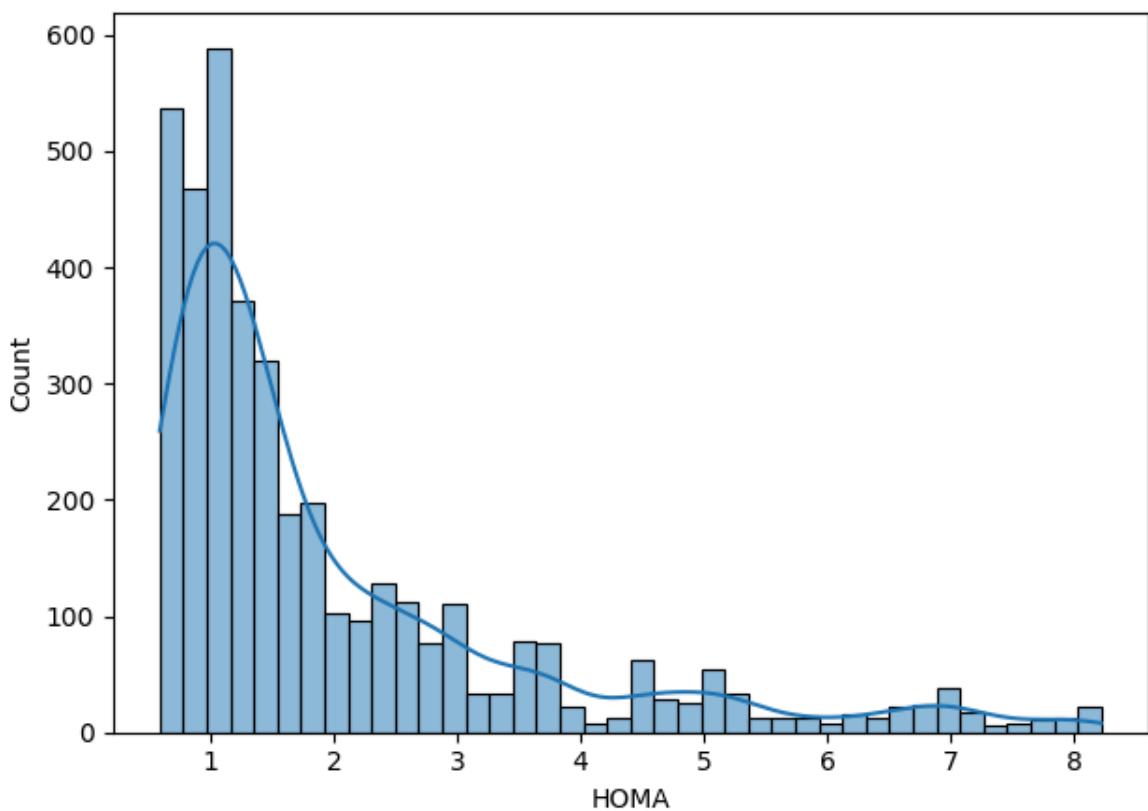
Distribution of Glucose



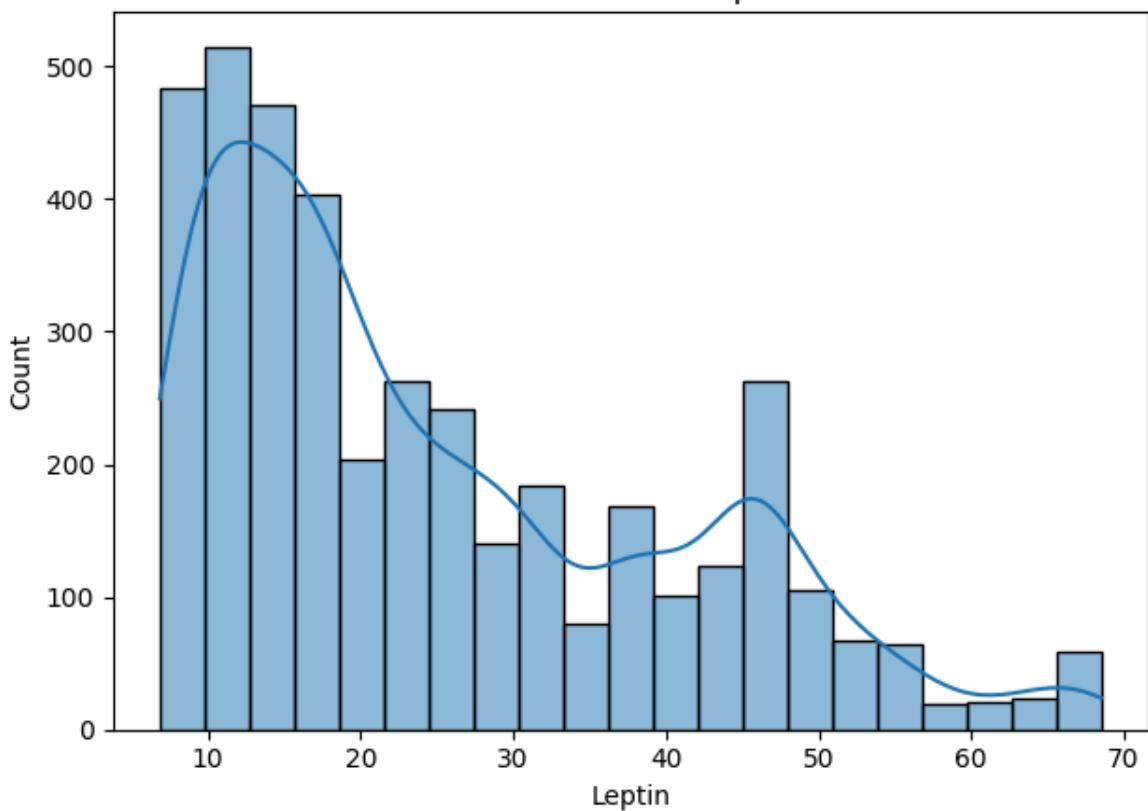
Distribution of Insulin

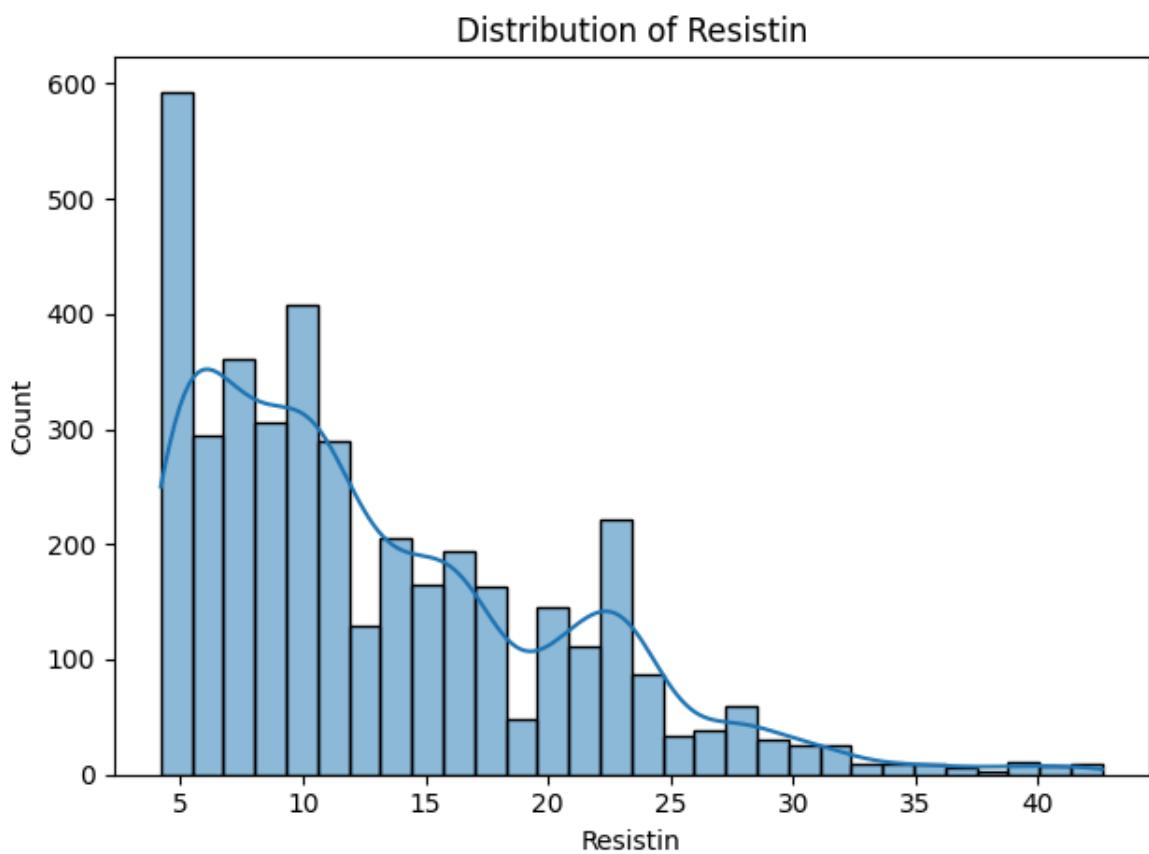
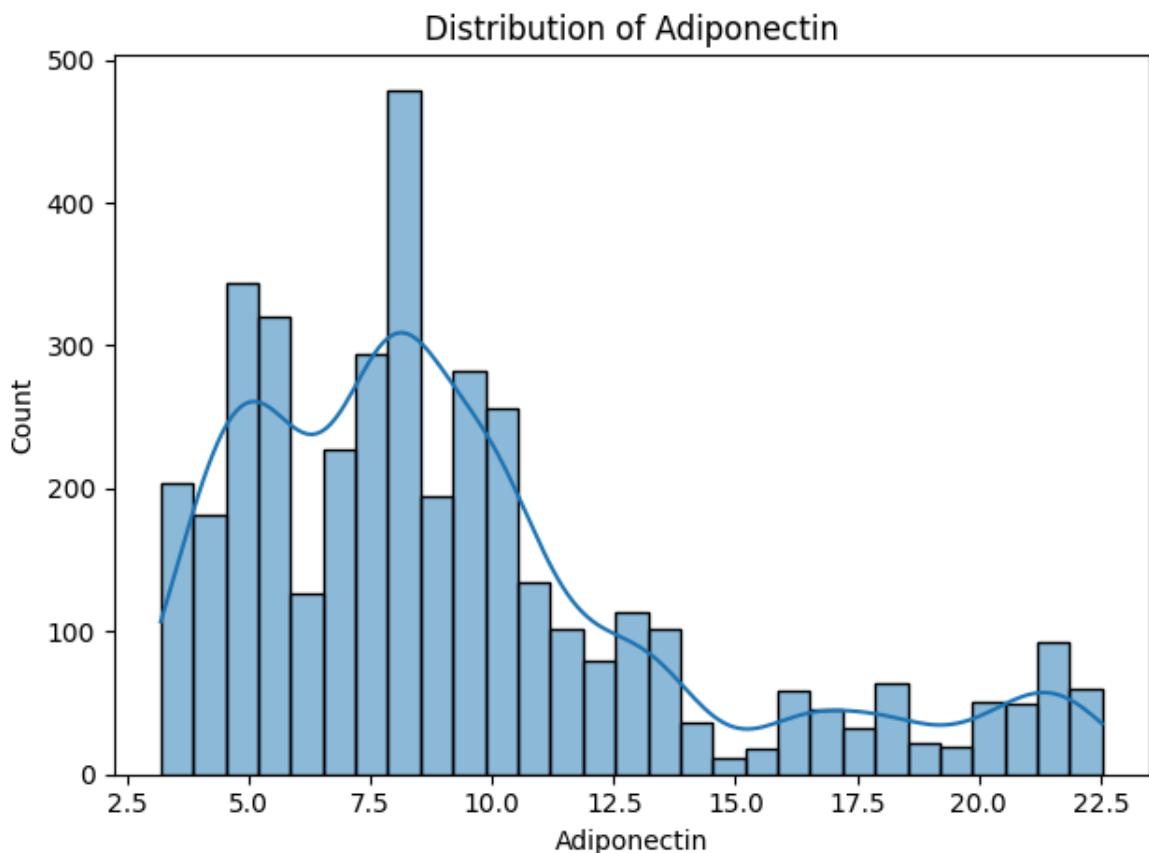


Distribution of HOMA

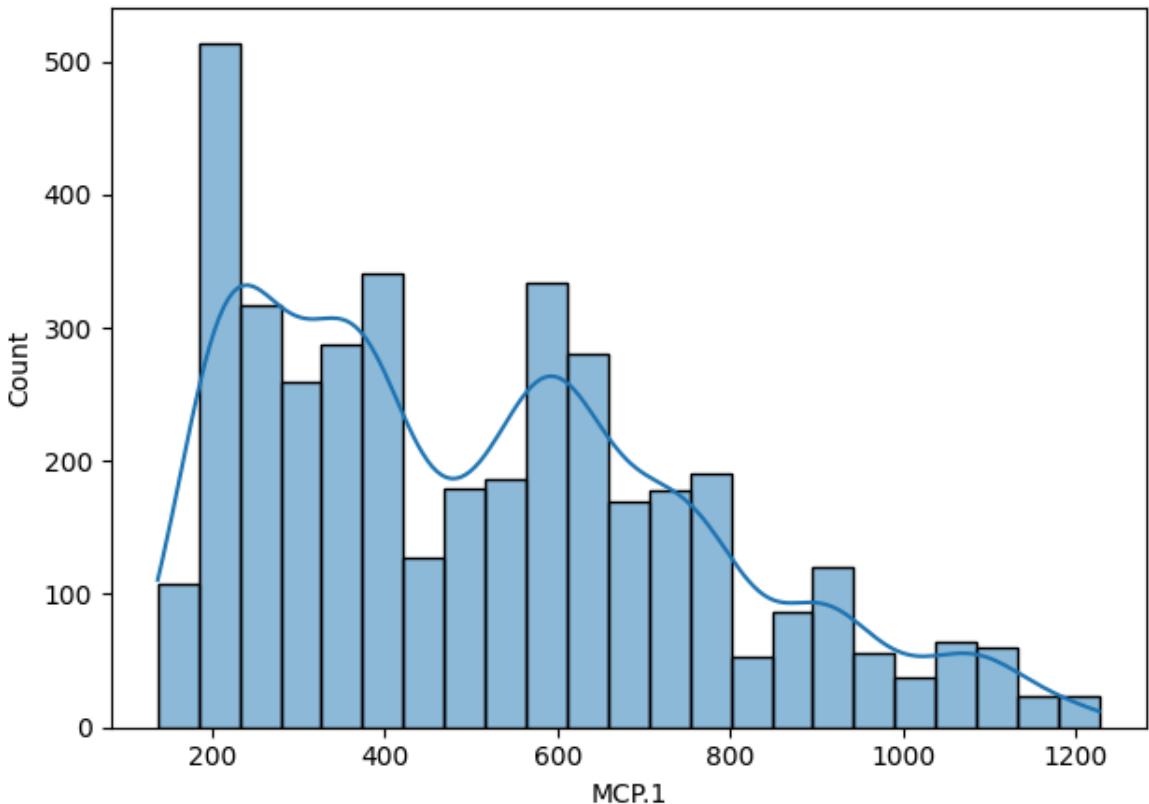


Distribution of Leptin





Distribution of MCP.1



3. Distributions by Class (Healthy vs Patient)

Overlaid histograms show how each feature differs between healthy and patient groups.

Key Observations

- Most features show **substantial overlap** between classes.
- **Insulin**, **HOMA**, and **Resistin** show slightly higher values for patients, but not strongly separated.
- **Glucose** distributions are similar in shape for both groups.

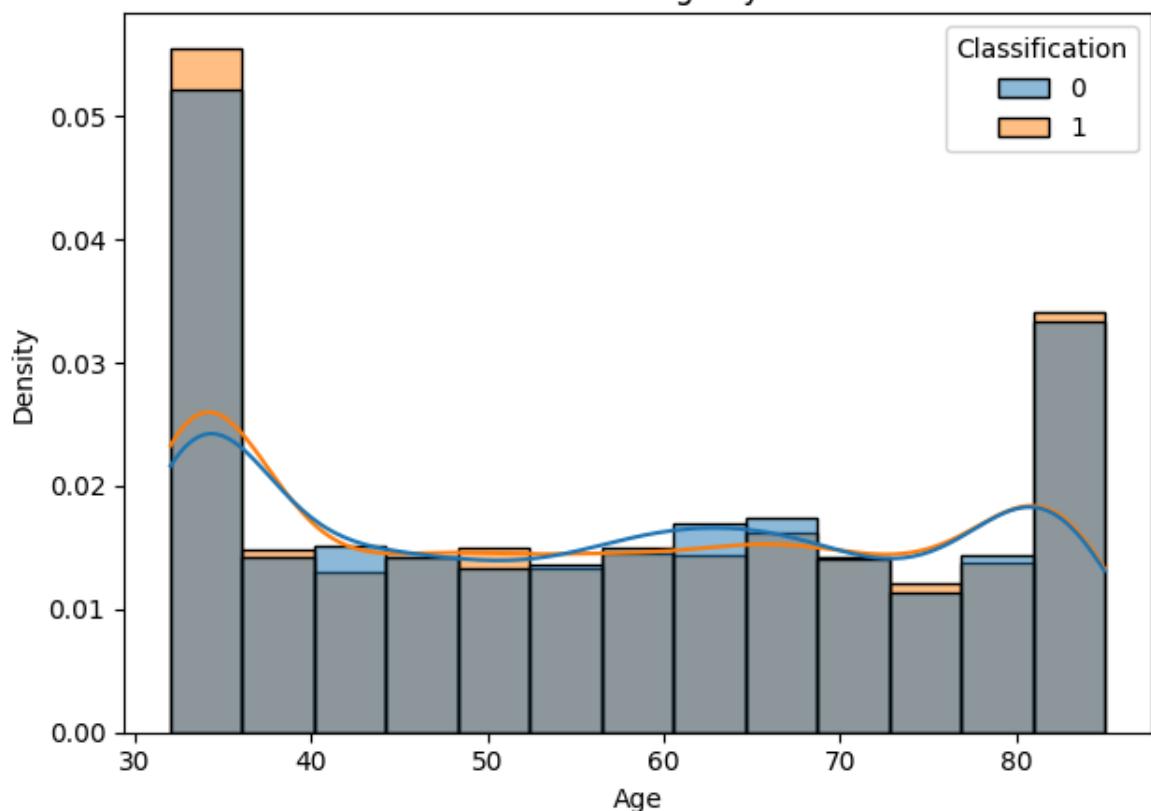
Conclusion

No single metabolic feature exhibits strong class-separation.

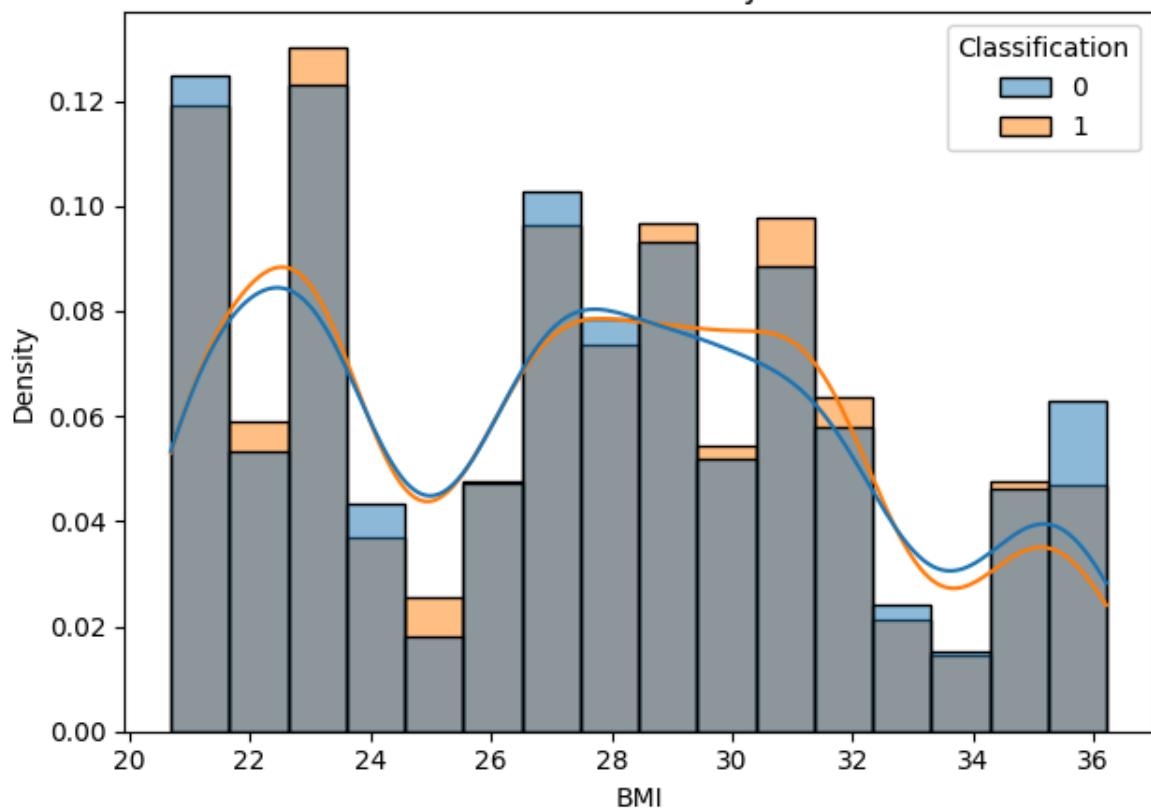
This suggests the need for **multivariate** and **nonlinear** modeling strategies.

```
In [13]: # Histograms by class (overlaid)
for col in feature_cols:
    fig, ax = plt.subplots()
    sns.histplot(data=df, x=col, hue=TARGET_COL, kde=True, stat="density", common_norm=False)
    ax.set_title(f"Distribution of {col} by Class")
    ax.set_xlabel(col)
    plt.tight_layout()
    plt.show()
```

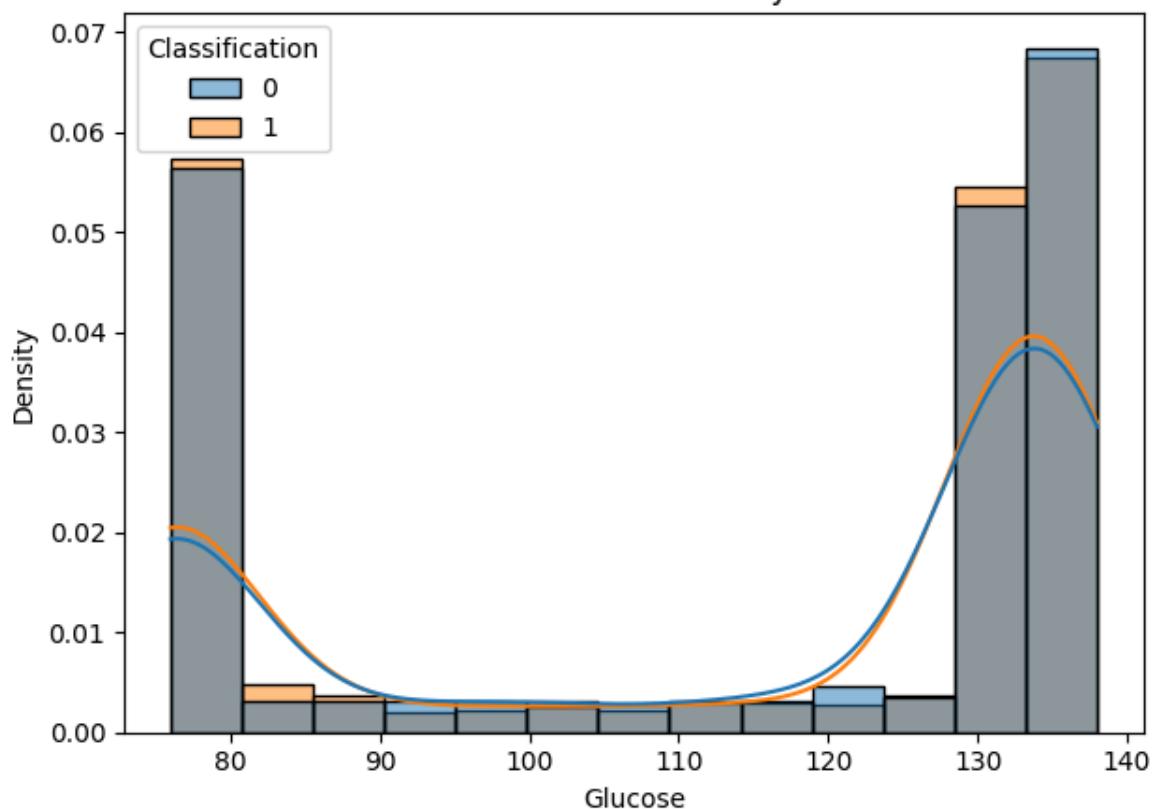
Distribution of Age by Class



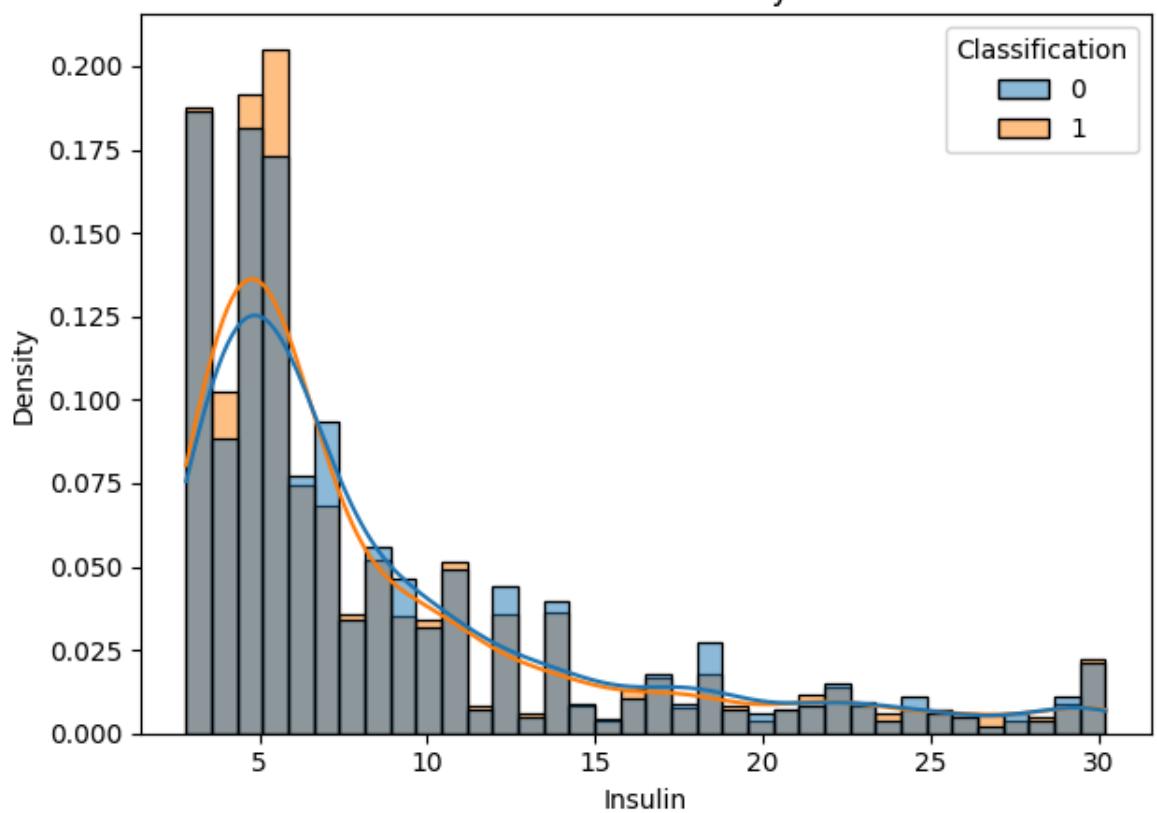
Distribution of BMI by Class



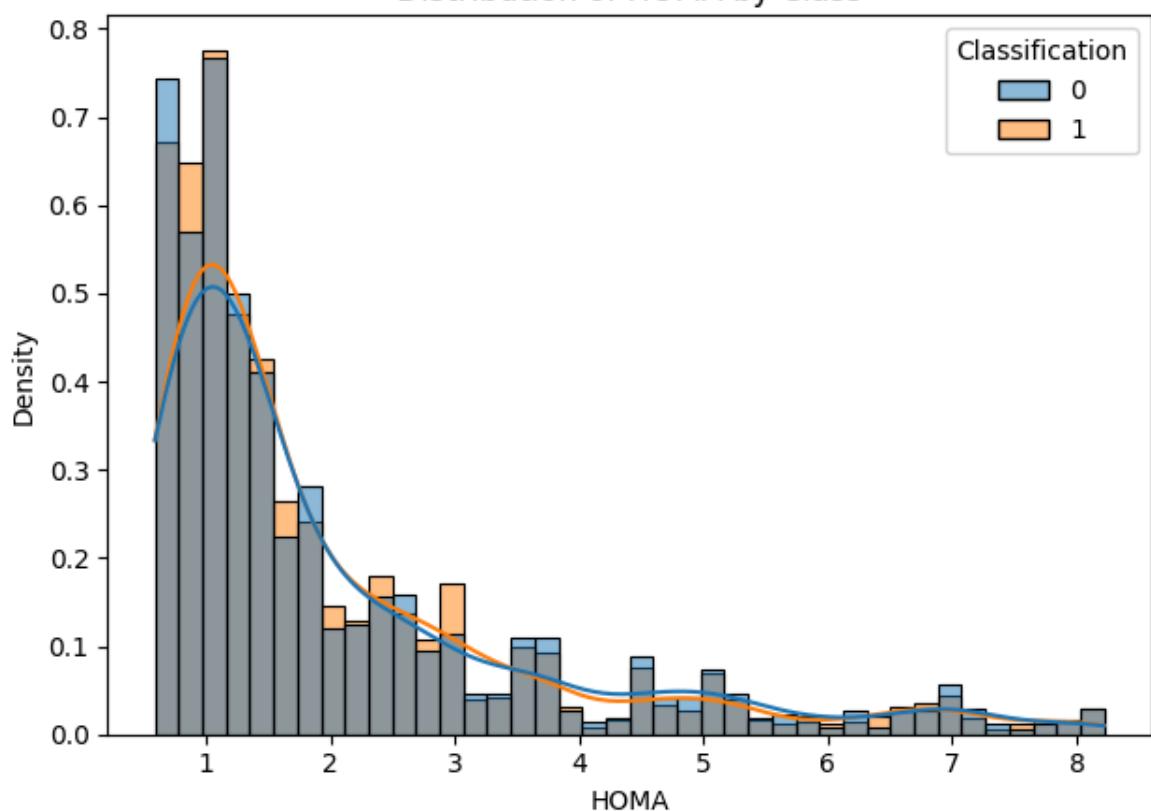
Distribution of Glucose by Class



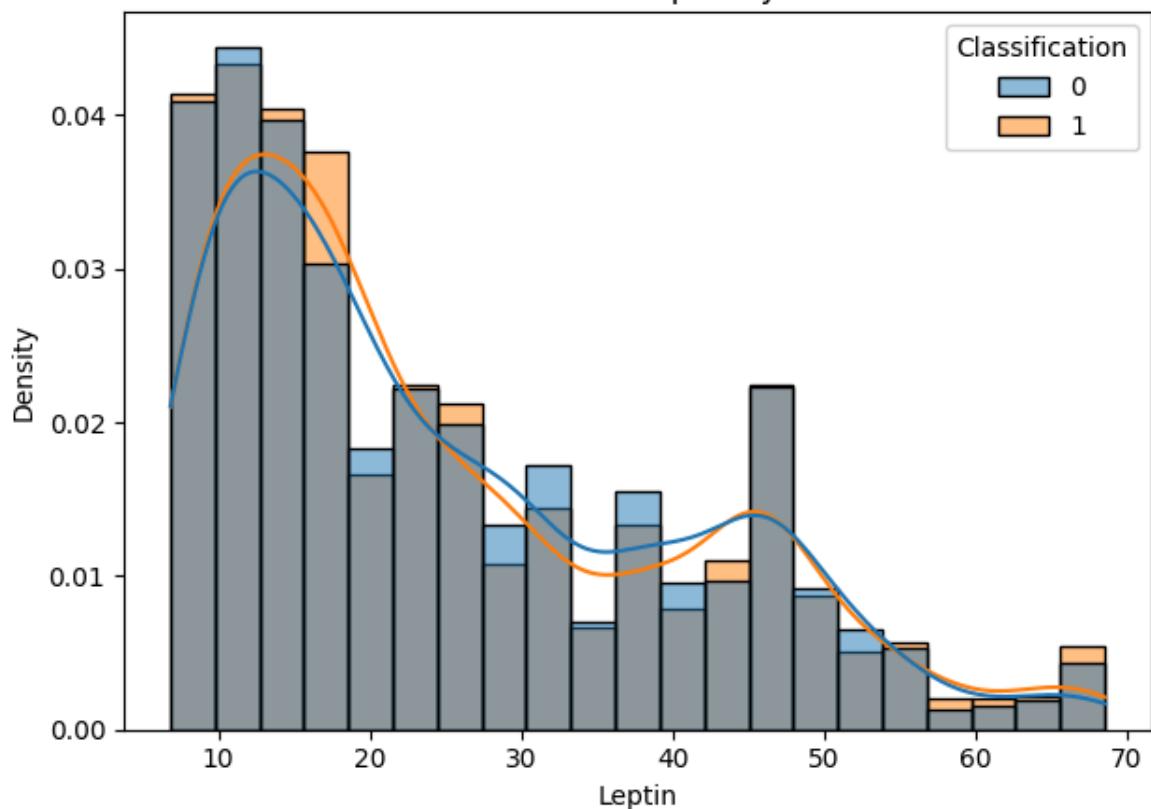
Distribution of Insulin by Class



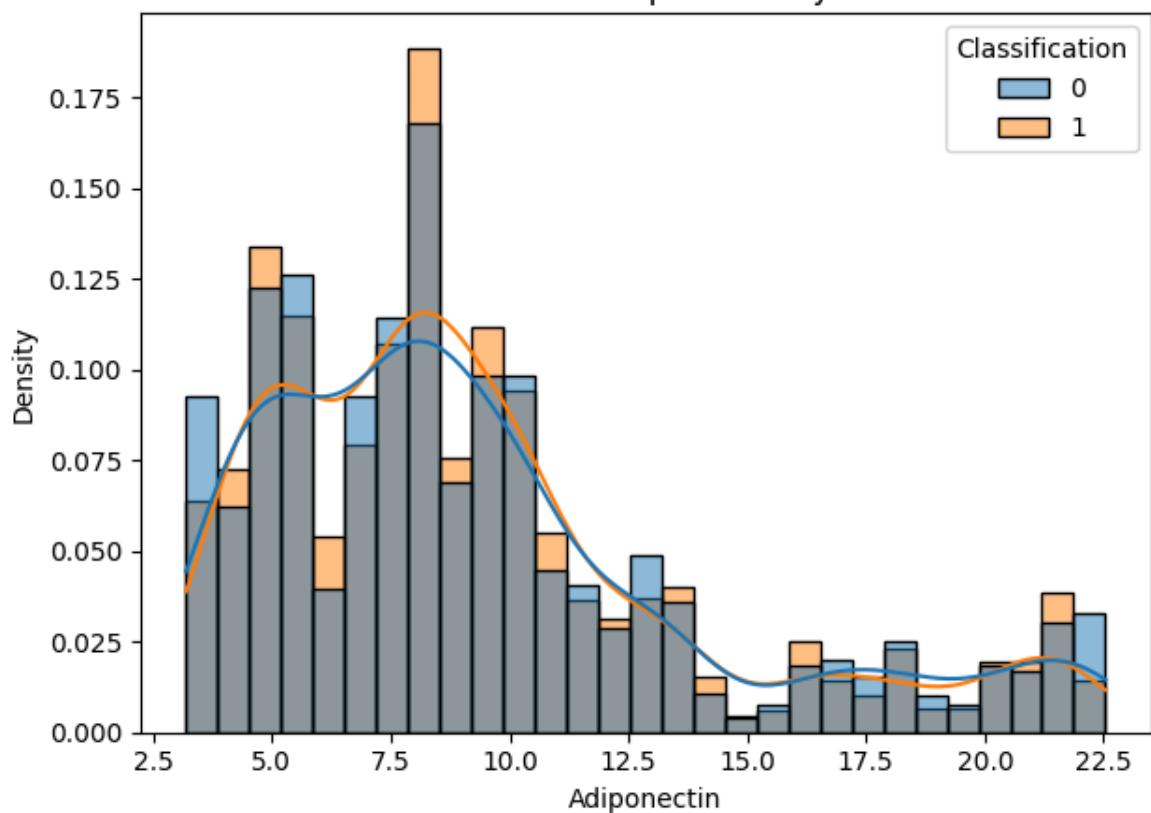
Distribution of HOMA by Class



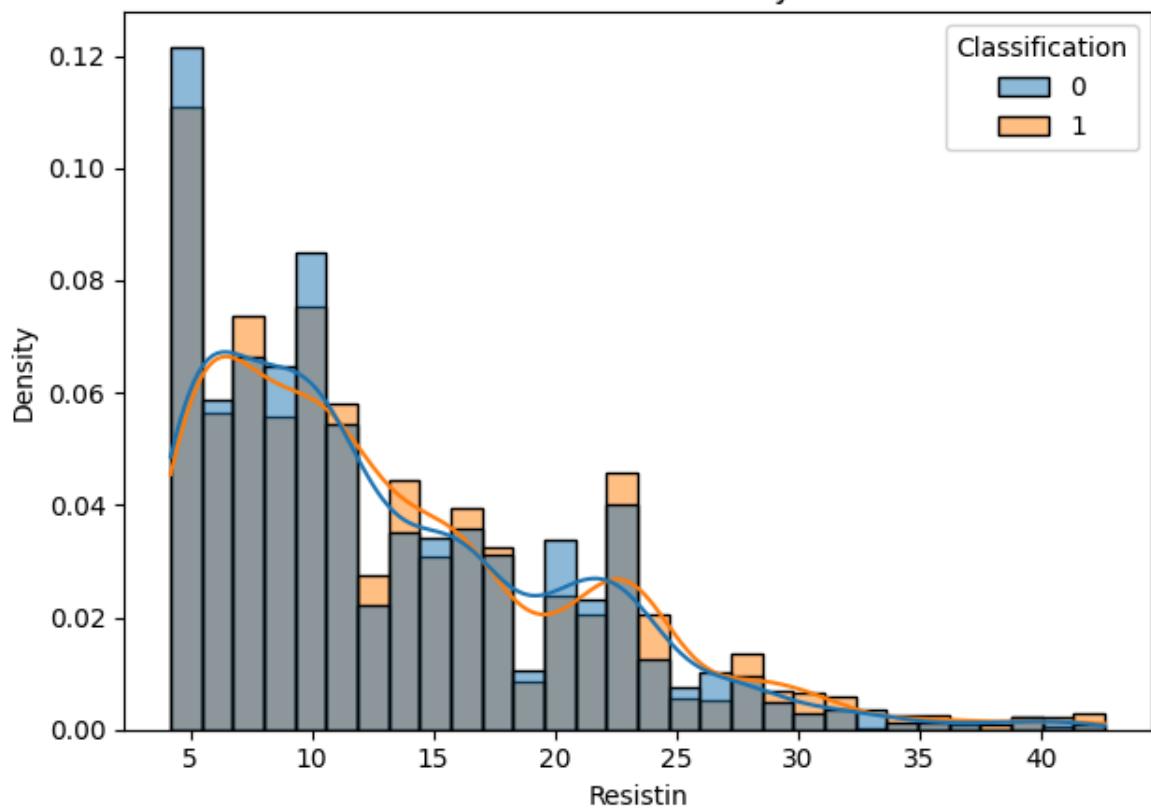
Distribution of Leptin by Class



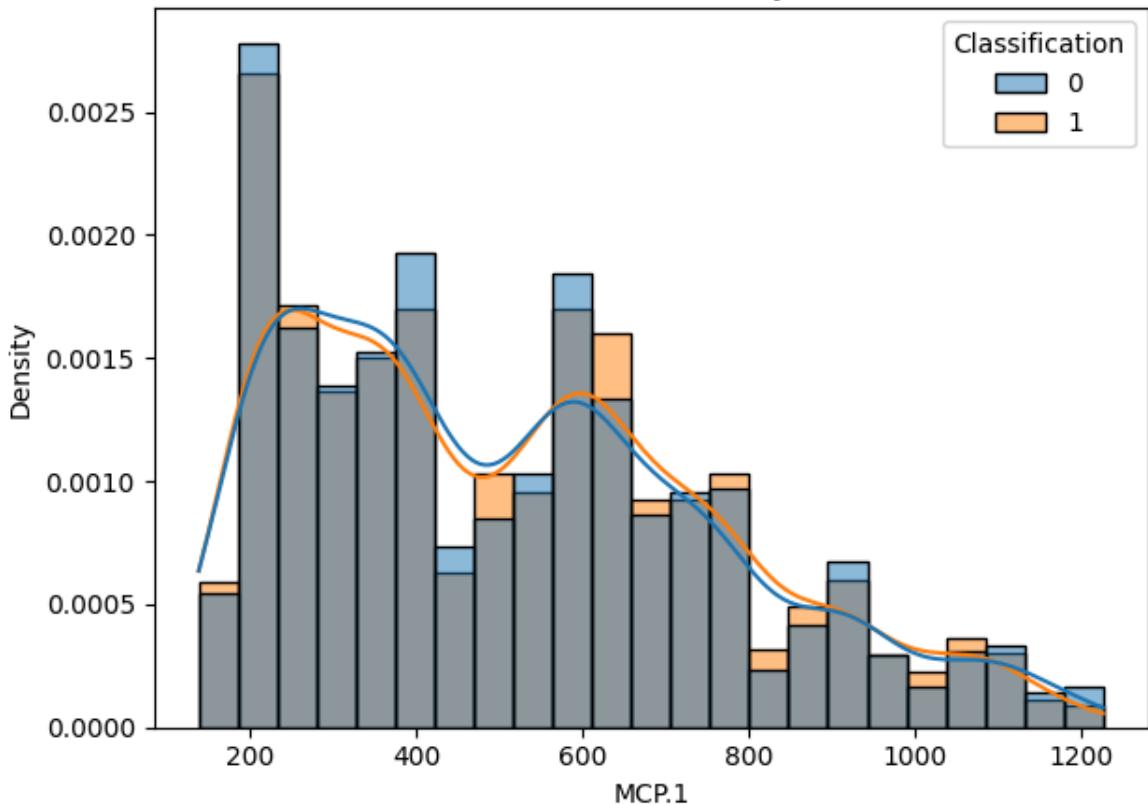
Distribution of Adiponectin by Class



Distribution of Resistin by Class



Distribution of MCP.1 by Class



4. Boxplots & Violin Plots by Class

Boxplots and violin plots confirm the overlapping distributions between healthy and patient groups.

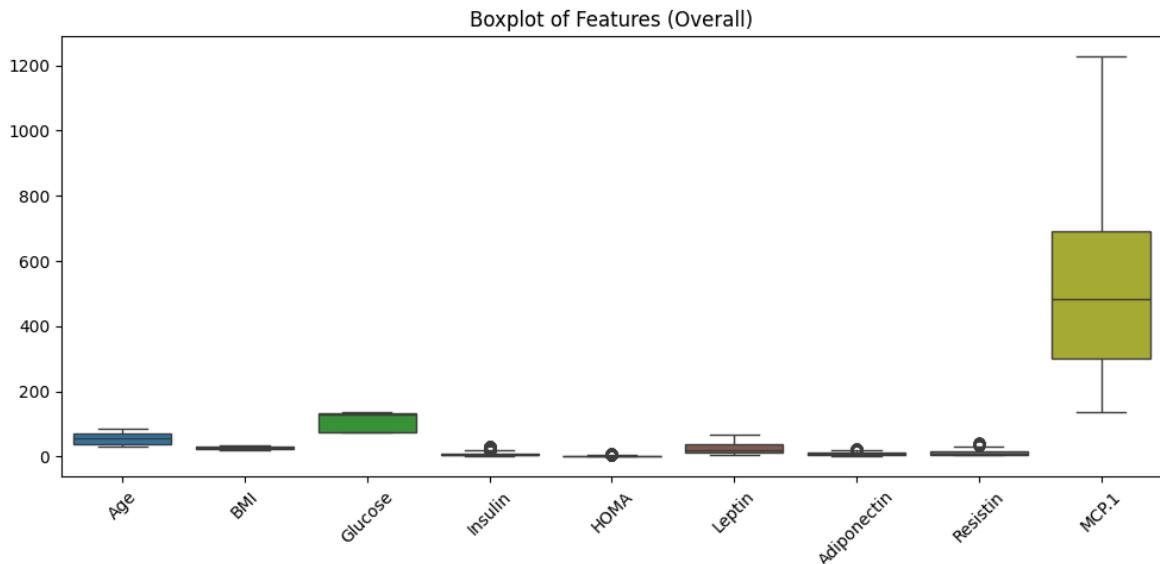
Observations

- Medians for all features are nearly identical between classes.
- Significant outliers are present in:
 - Insulin
 - HOMA
 - Resistin
- The violin plot shapes further reinforce the lack of clear separation.

Interpretation

The dataset likely contains **weak individual predictors**, requiring combined or interaction-based modeling.

```
In [14]: # Overall boxplots (not stratified by class)
plt.figure(figsize=(10, 5))
sns.boxplot(data=df[feature_cols])
plt.xticks(rotation=45)
plt.title("Boxplot of Features (Overall)")
plt.tight_layout()
plt.show()
```



```
In [15]: # Summary statistics per class
grouped_stats = df.groupby(TARGET_COL)[feature_cols].describe().transpose()
grouped_stats
```

Out[15]:

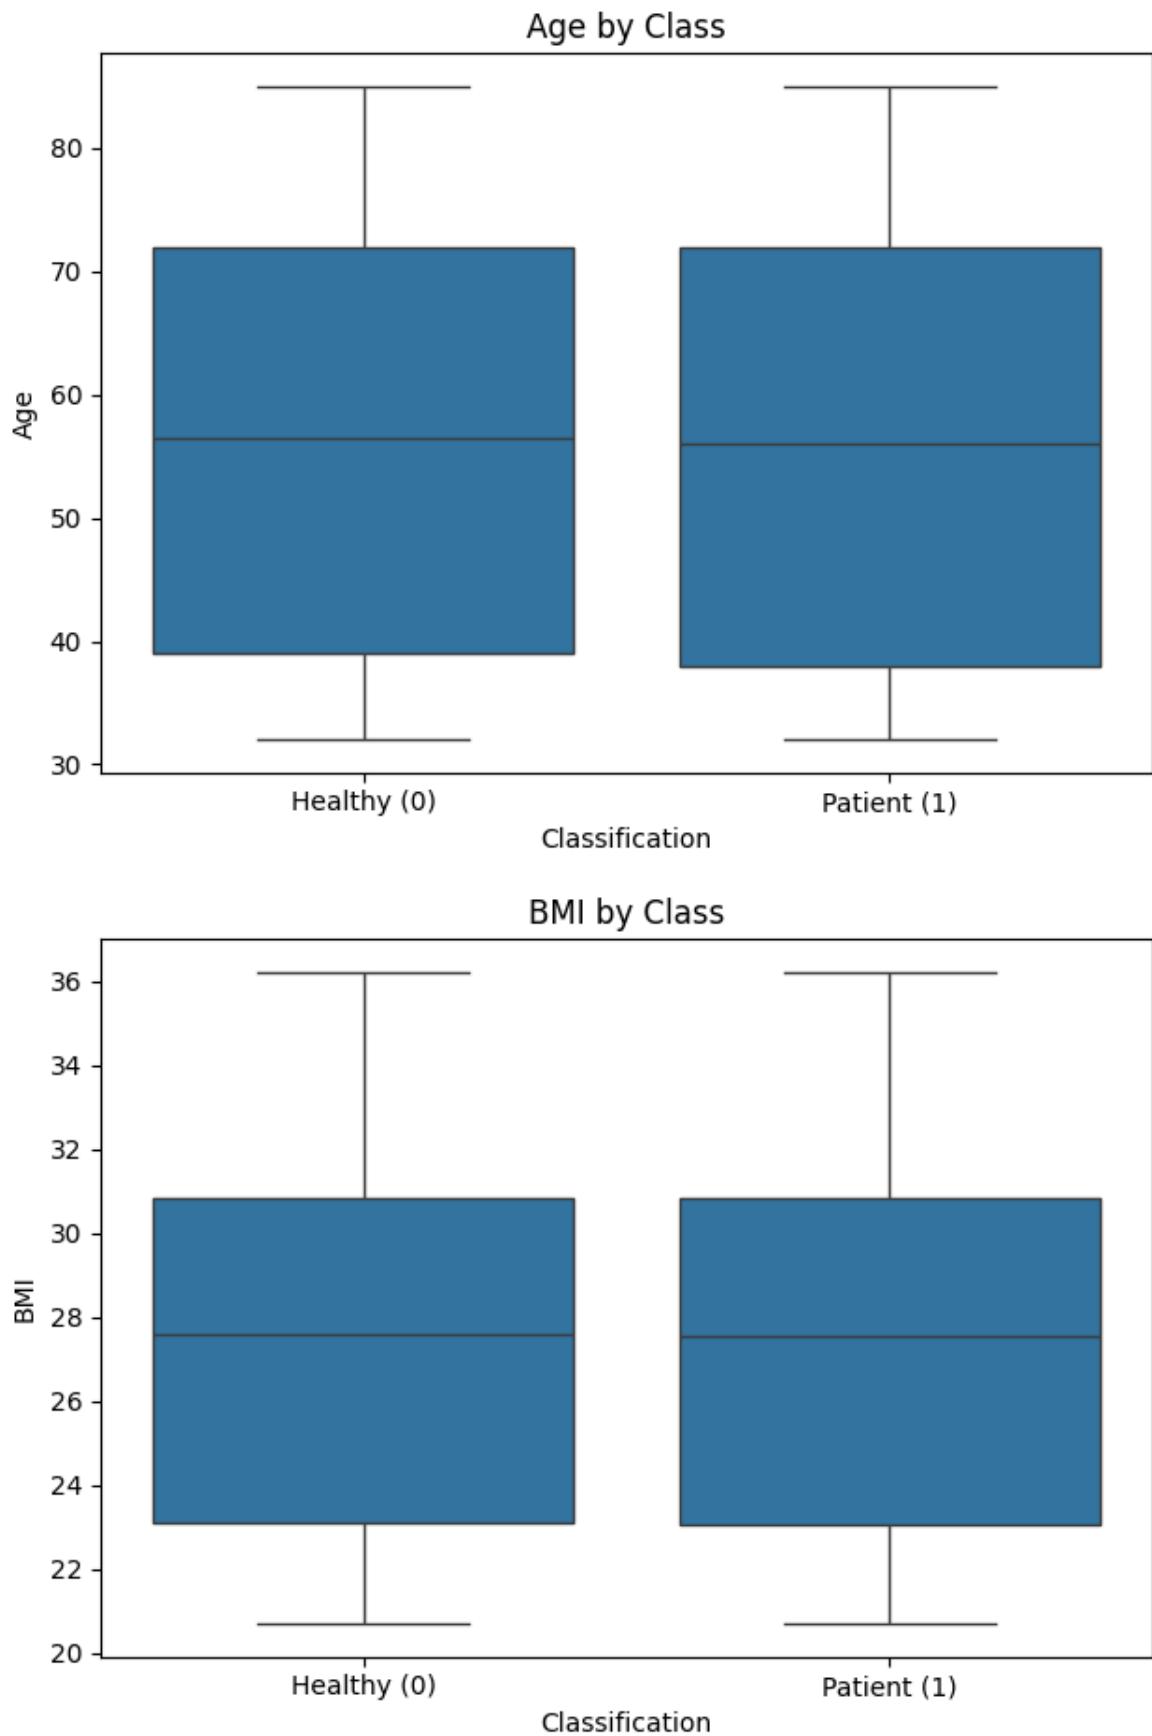
	Classification	0	1
Age	count	1784.000000	2216.000000
	mean	56.368834	56.083484
	std	17.702713	17.898252
	min	32.000000	32.000000
	25%	39.000000	38.000000
...	
MCP.1	min	137.488000	138.903000
	25%	300.156250	299.263750
	50%	473.118000	483.763000
	75%	679.262500	696.078500
	max	1223.677000	1227.250000

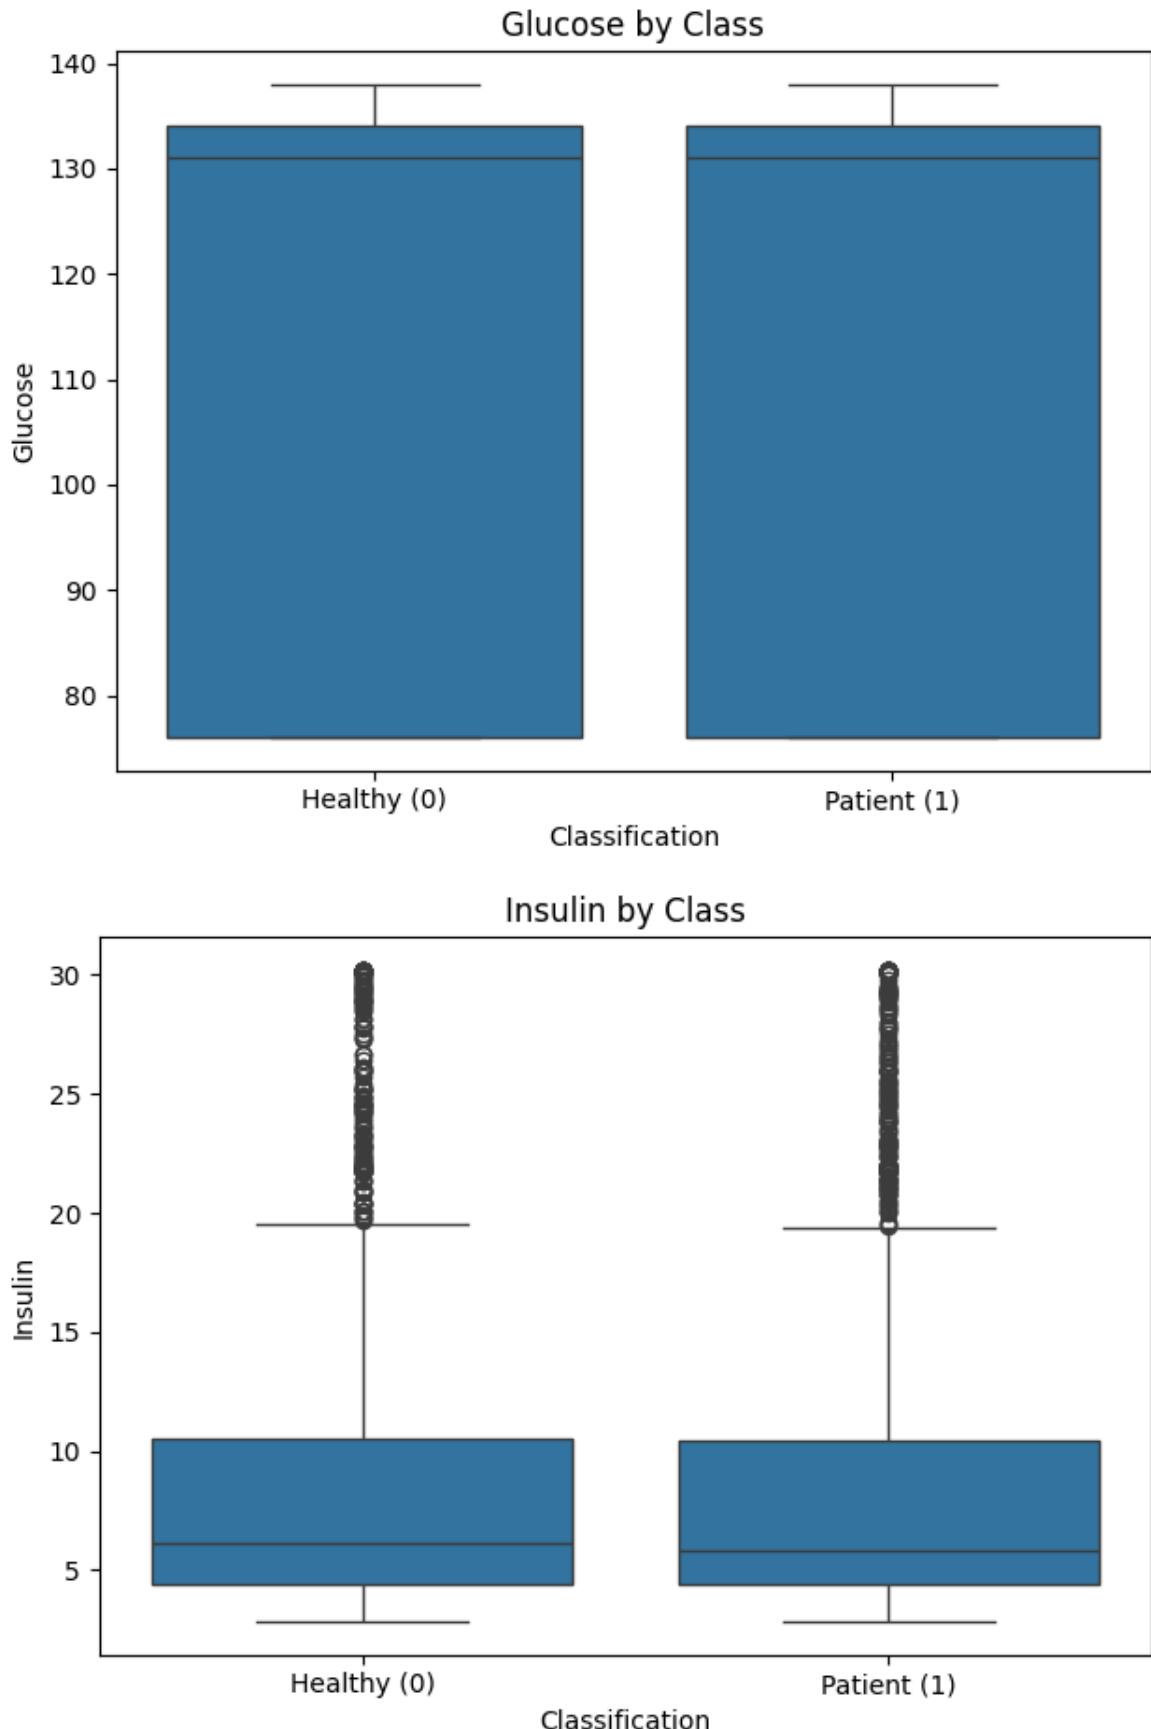
72 rows × 2 columns

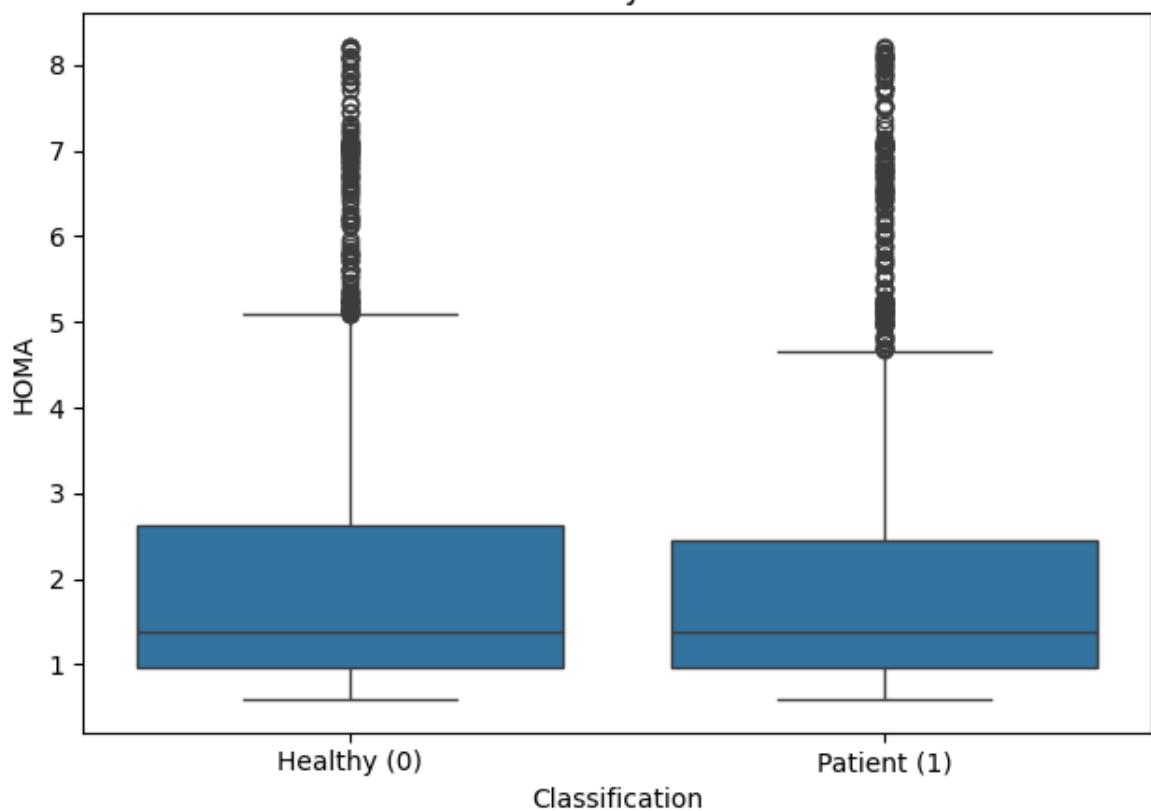
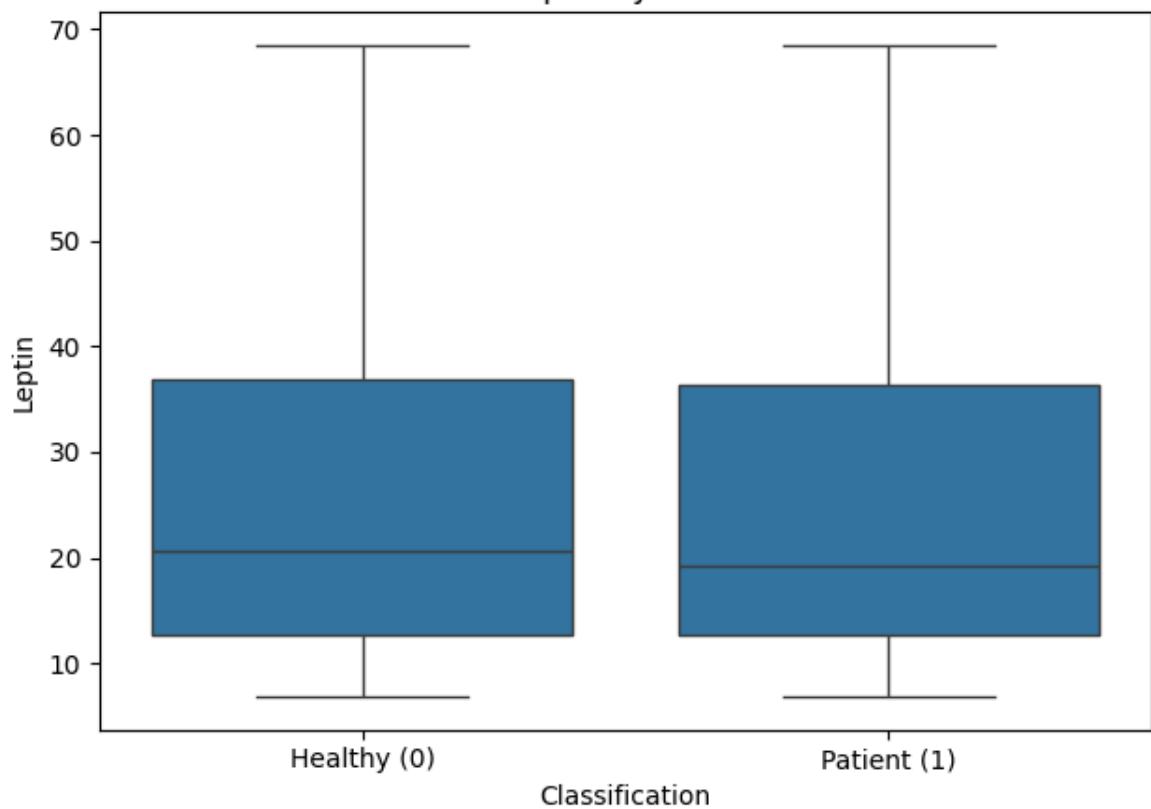
```
In [ ]: for col in feature_cols:
    fig, ax = plt.subplots()
    sns.boxplot(x=TARGET_COL, y=col, data=df, ax=ax)

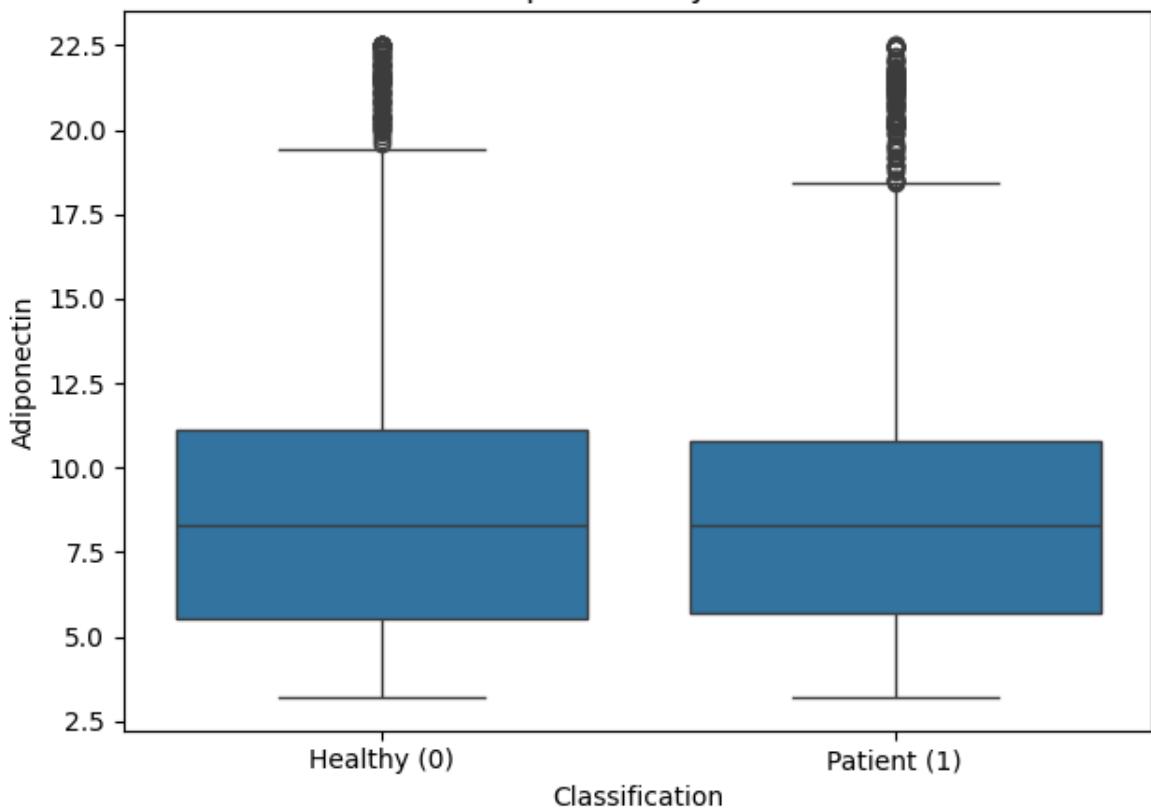
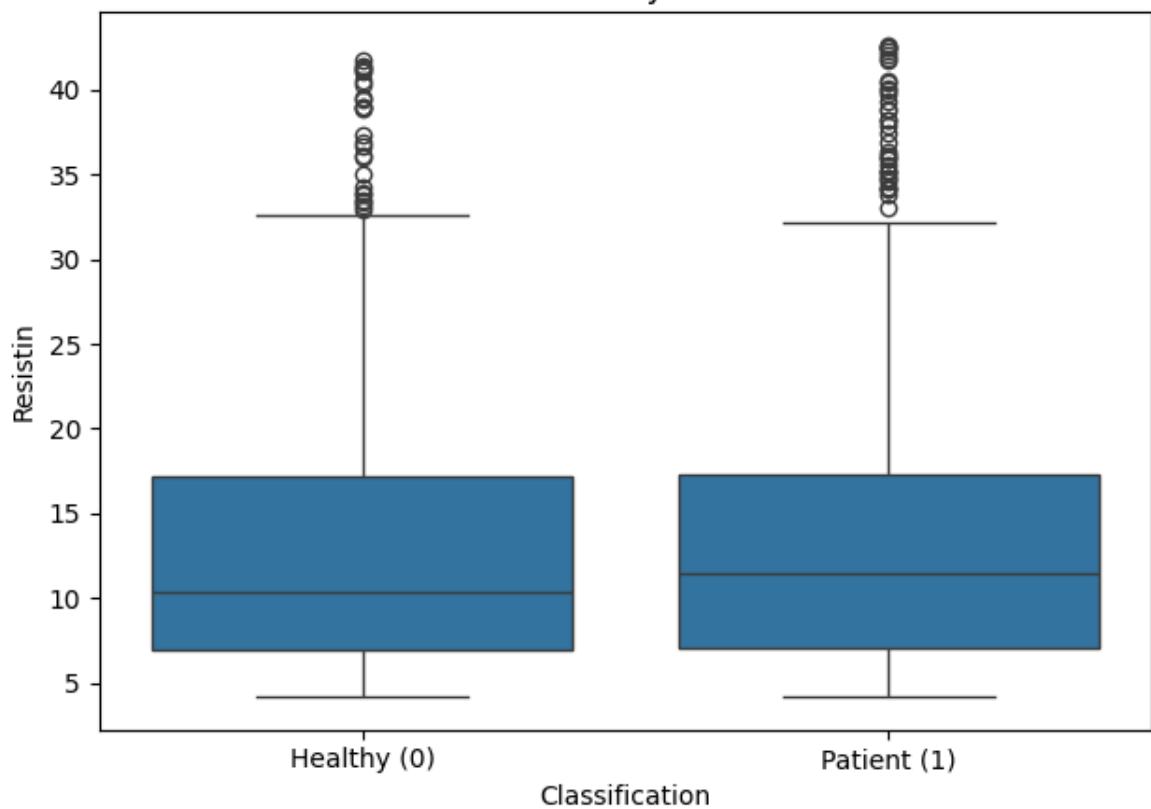
    ax.set_xticks([0, 1])
    ax.set_xticklabels(["Healthy (0)", "Patient (1)"])

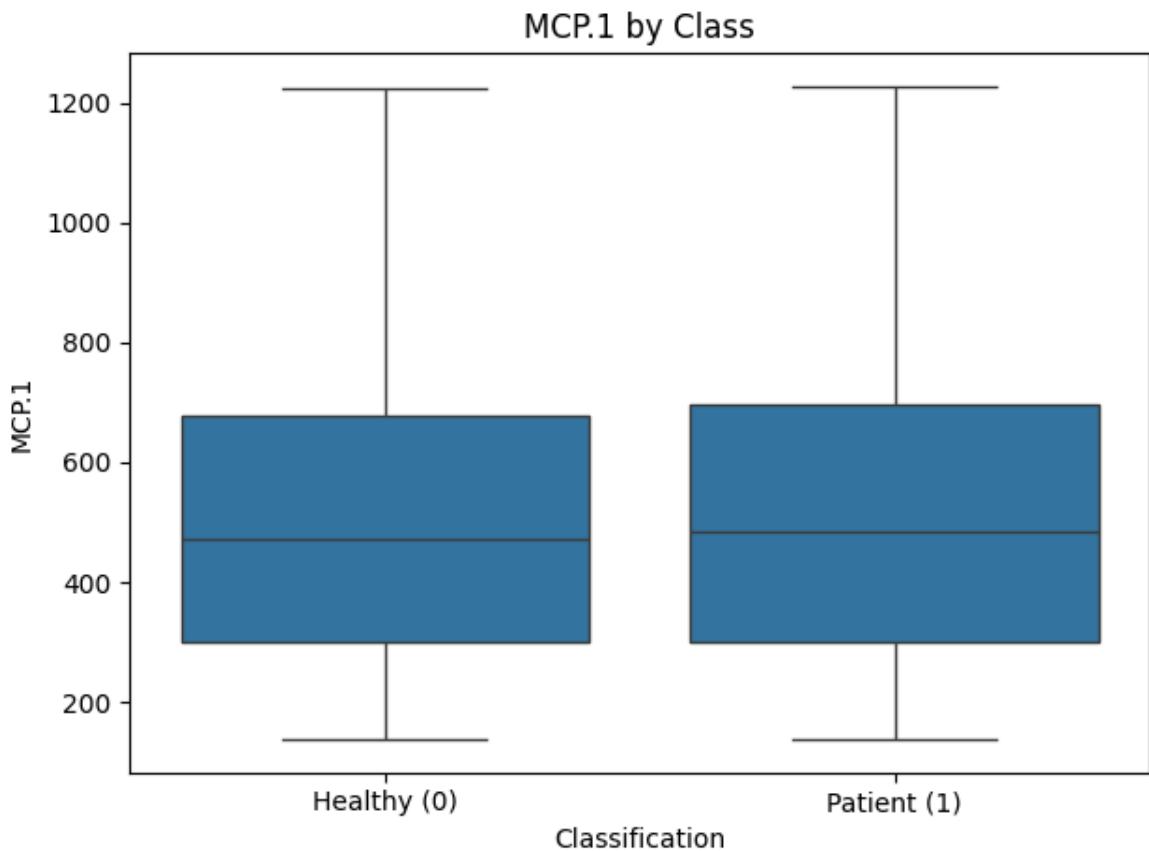
    ax.set_title(f"{col} by Class")
    plt.tight_layout()
    plt.show()
```





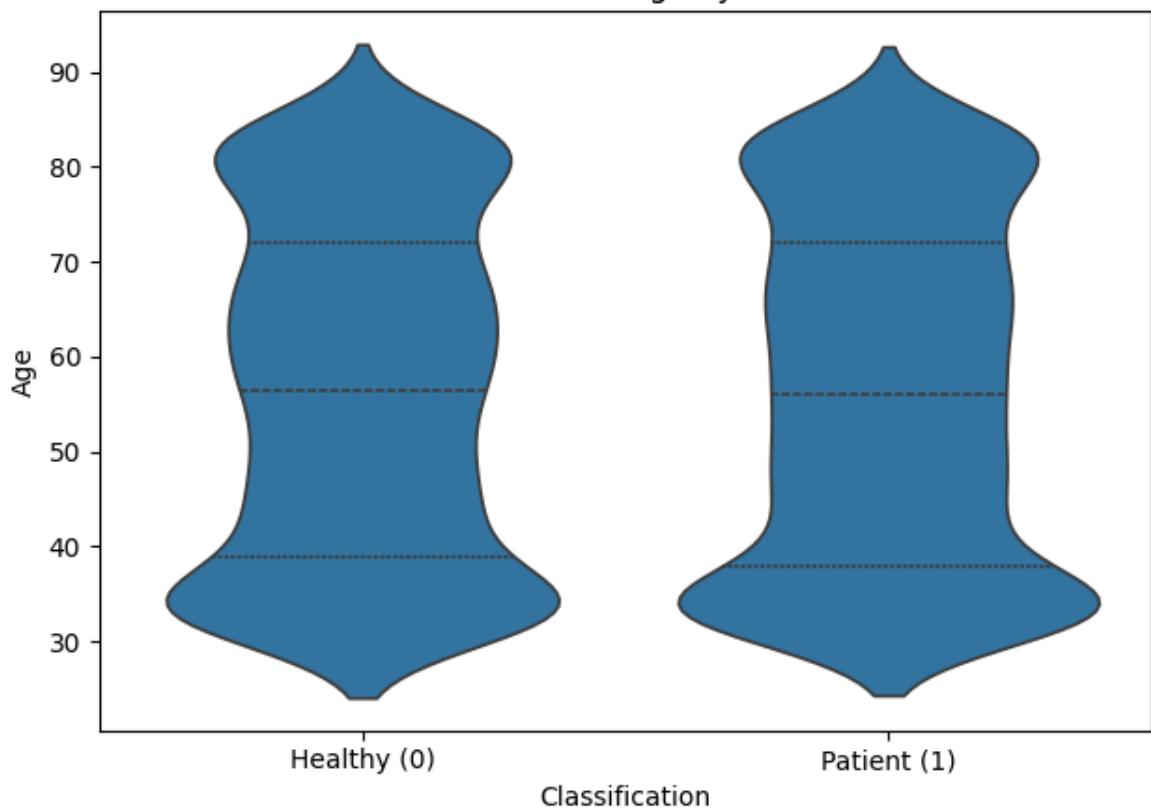
HOMA by Class**Leptin by Class**

Adiponectin by Class**Resistin by Class**

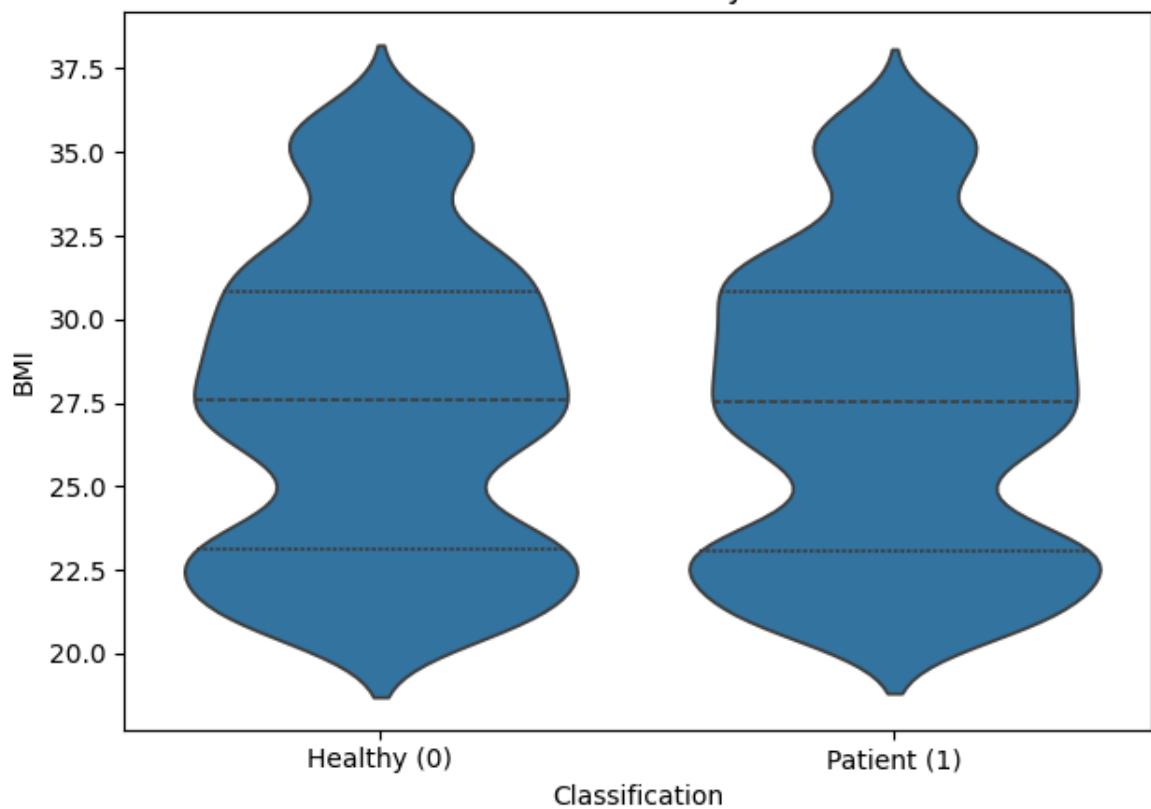


```
In [21]: for col in feature_cols:  
    fig, ax = plt.subplots()  
    sns.violinplot(x=TARGET_COL, y=col, data=df, inner="quartile", ax=ax)  
    ax.set_xticks([0, 1])  
    ax.set_xticklabels(["Healthy (0)", "Patient (1)"])  
    ax.set_title(f"Violin Plot of {col} by Class")  
    plt.tight_layout()  
    plt.show()
```

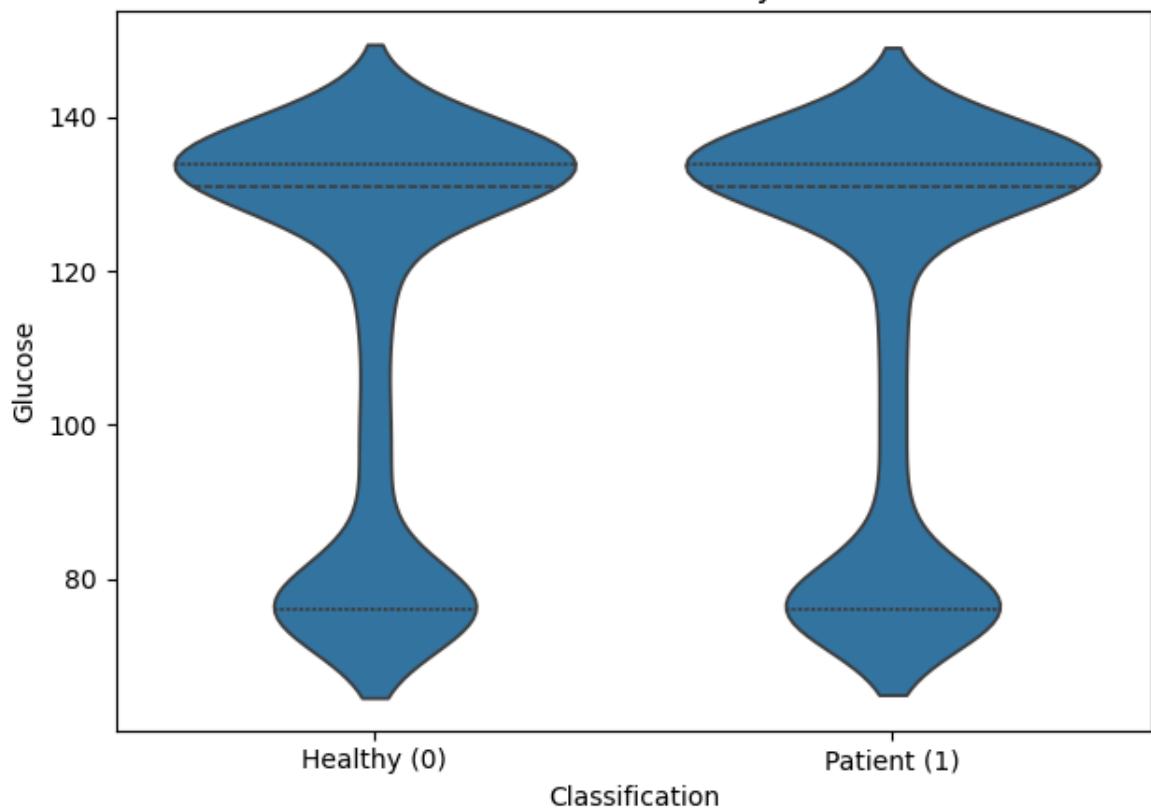
Violin Plot of Age by Class



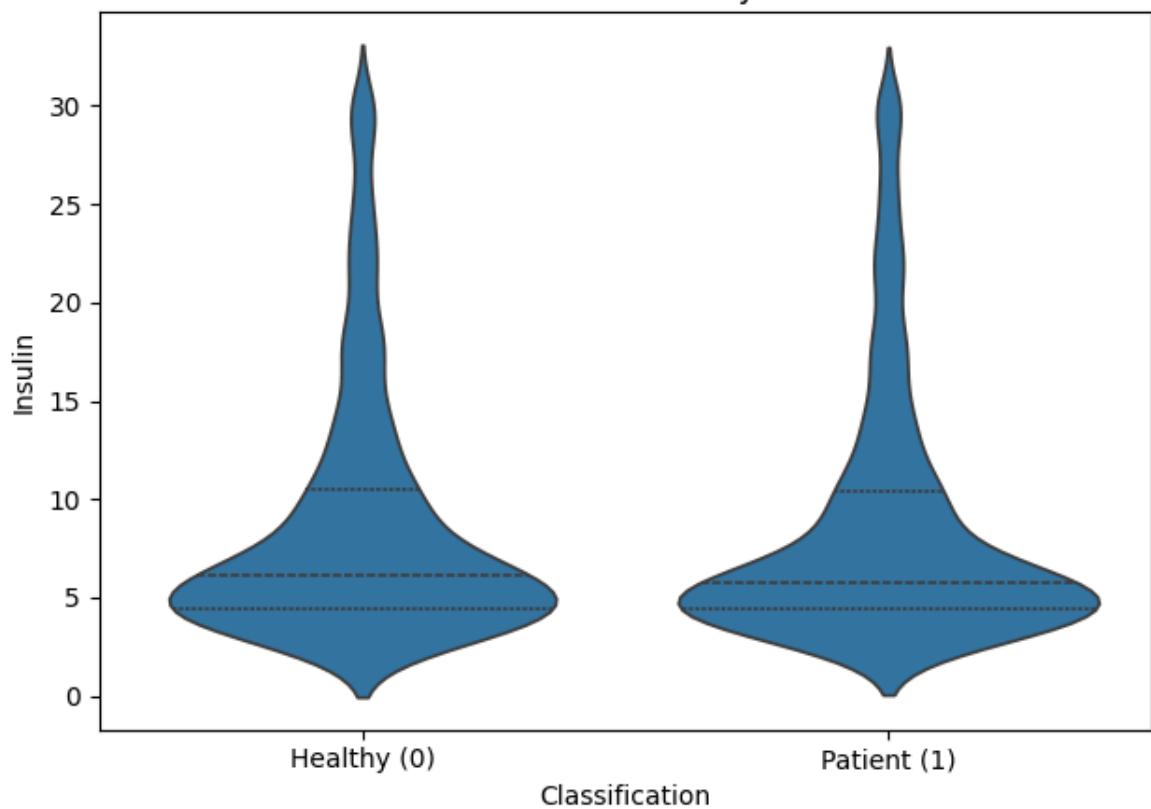
Violin Plot of BMI by Class



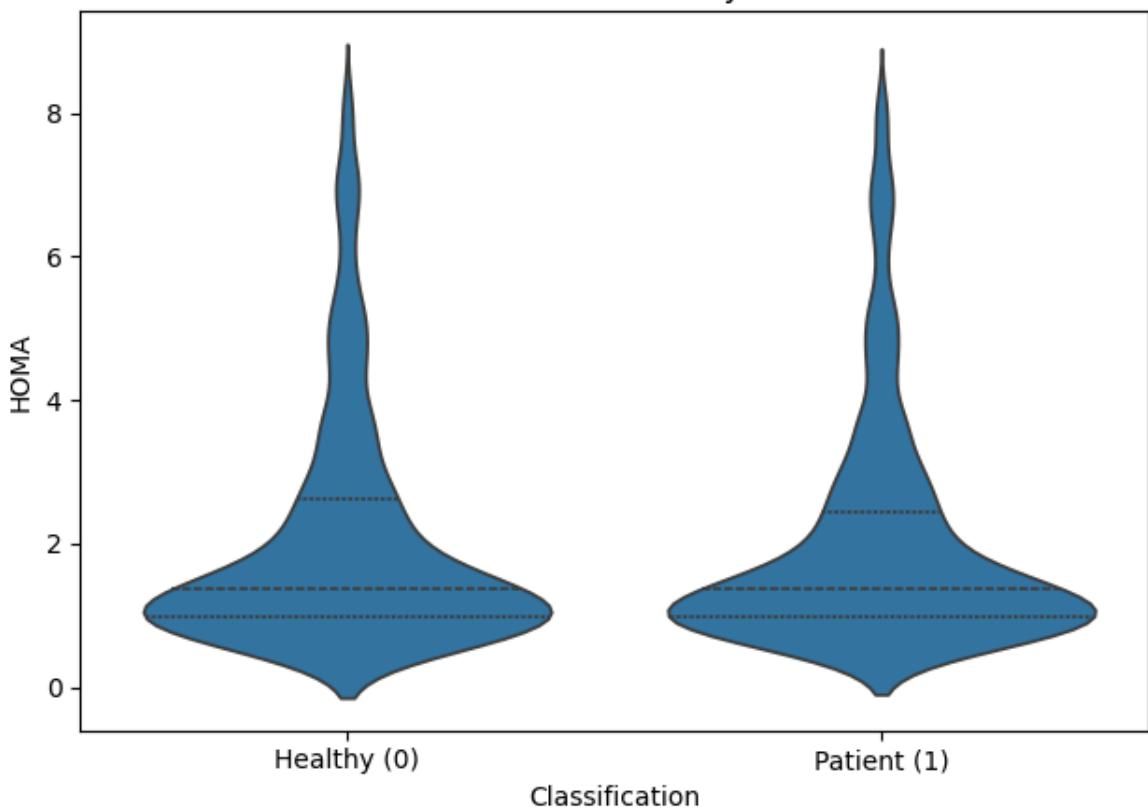
Violin Plot of Glucose by Class



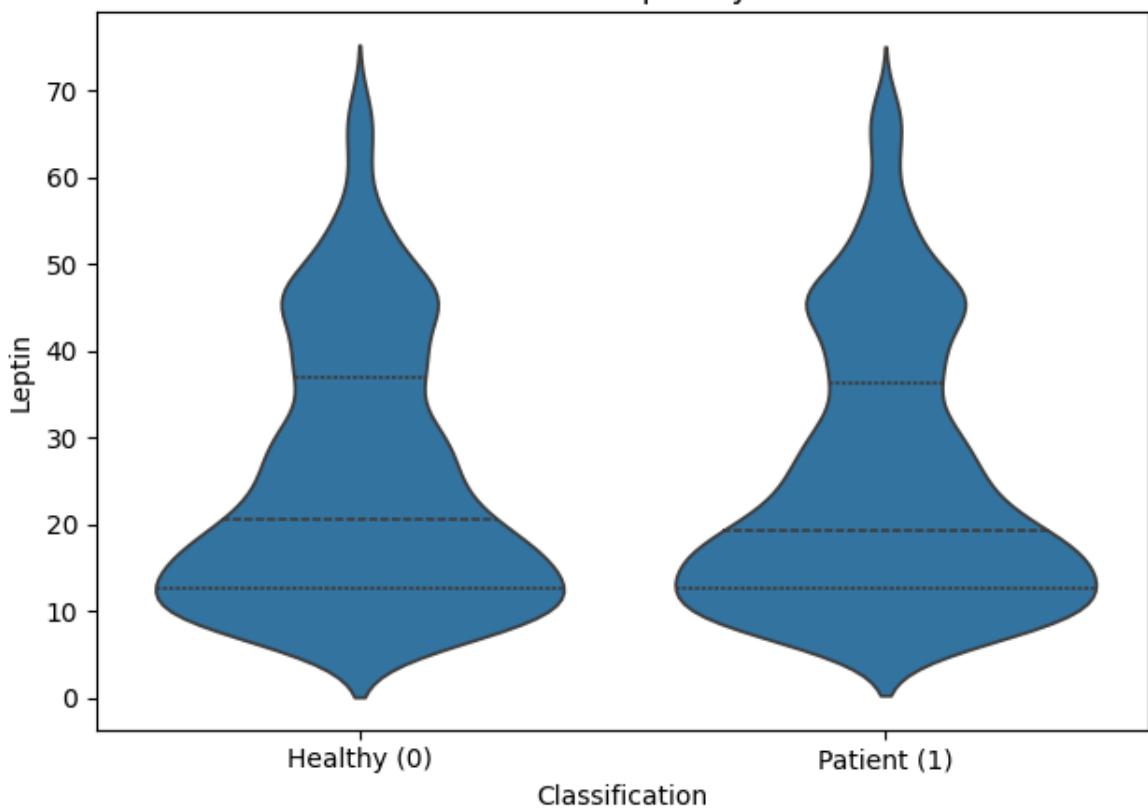
Violin Plot of Insulin by Class



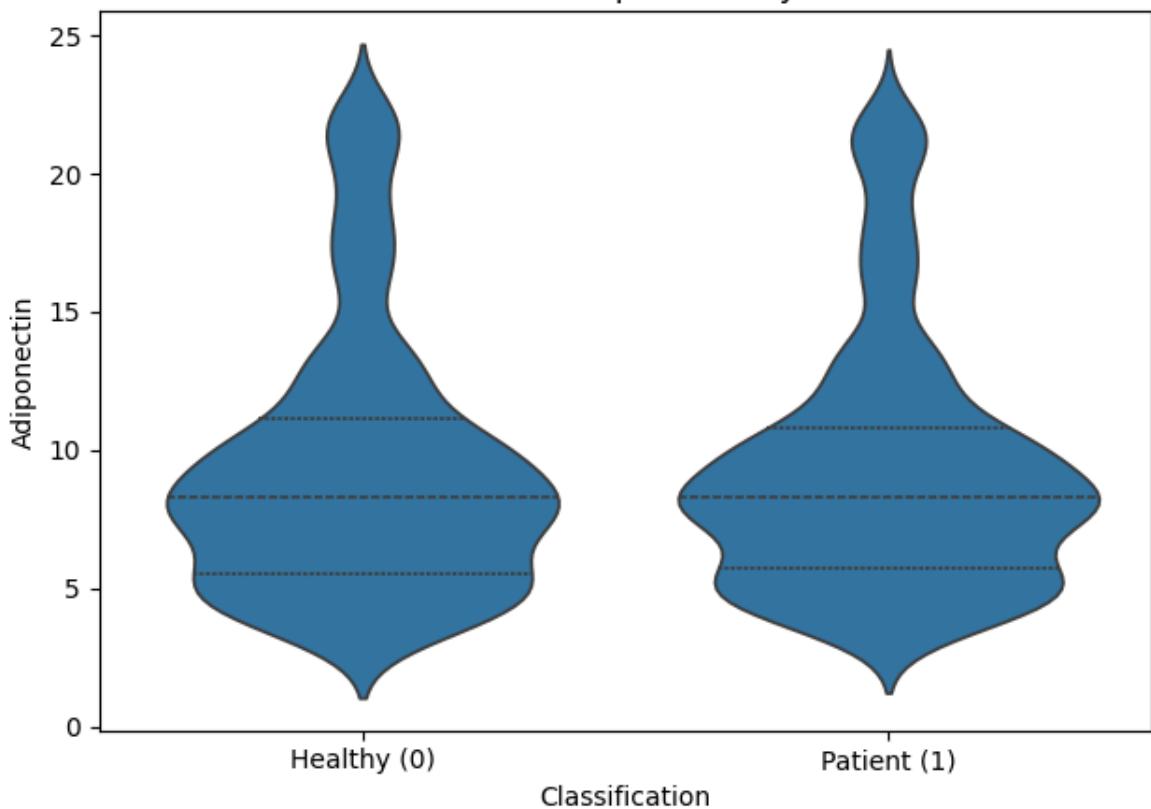
Violin Plot of HOMA by Class



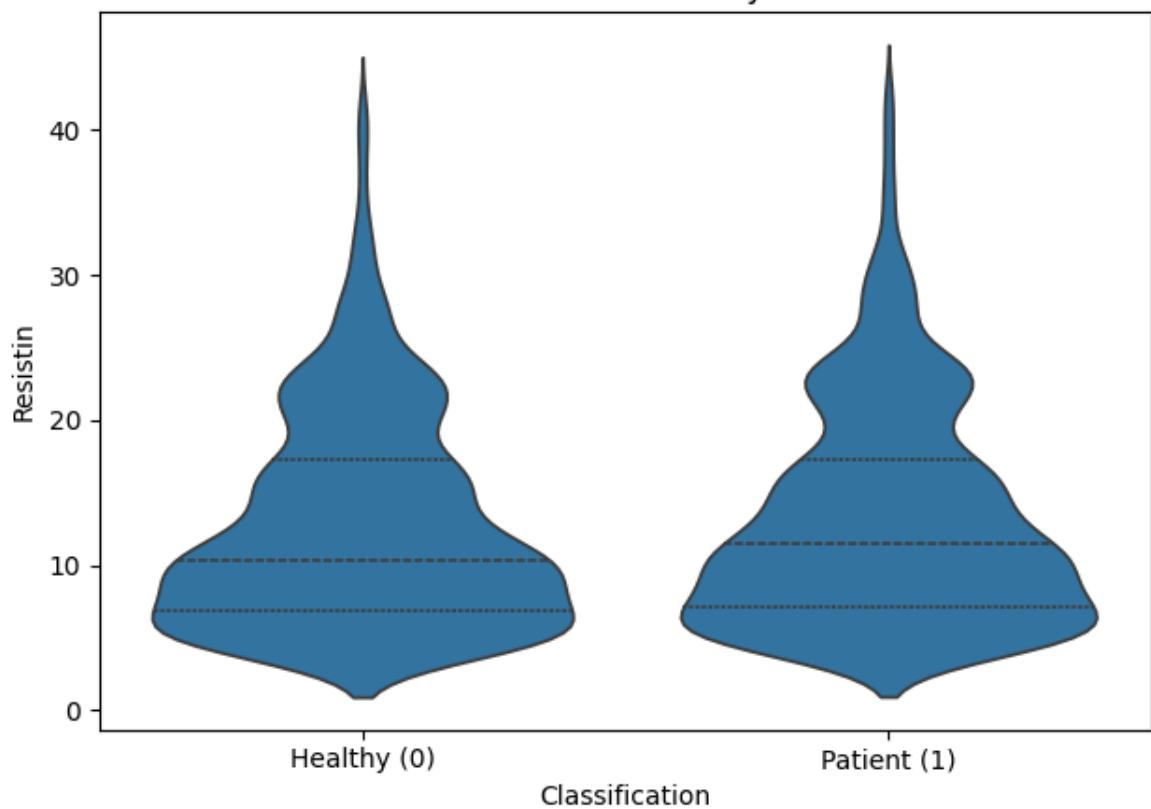
Violin Plot of Leptin by Class

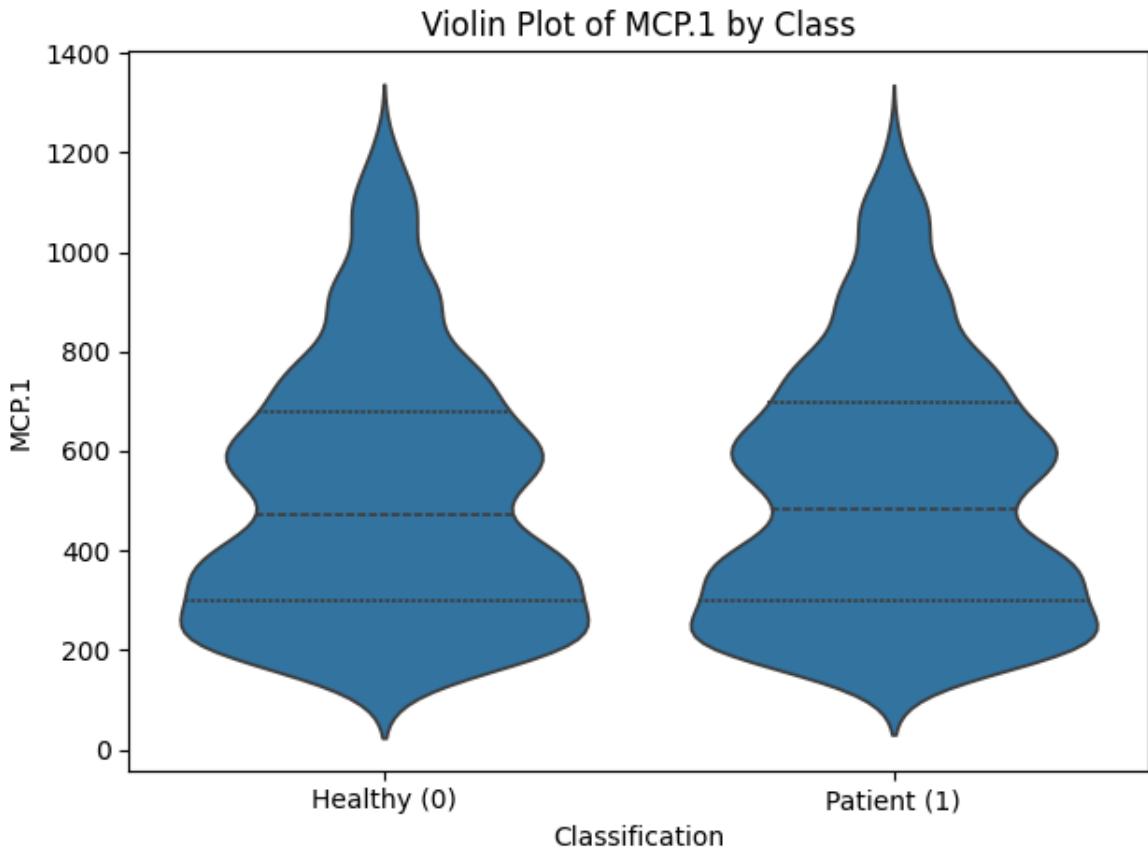


Violin Plot of Adiponectin by Class



Violin Plot of Resistin by Class





5. Statistical Tests (t-test & Mann-Whitney U)

For each feature, both **t-tests** and **Mann-Whitney U tests** were performed.

Most distinguishable features (still weak):

Feature	Mann-Whitney p-value	Interpretation
Insulin	0.074	Weak difference; borderline significance
Resistin	0.091	Weak difference
Others	> 0.6	No meaningful difference

Conclusion

- The dataset shows **very limited statistical separation** between classes.
- Insulin and Resistin contain minor signals but not strong enough alone.

```
In [22]: results = []

for col in feature_cols:
    # Split by class
    group0 = df[df[TARGET_COL] == 0][col]
    group1 = df[df[TARGET_COL] == 1][col]

    # Normality check (optional, small n so not super reliable)
    # Shapiro-Wilk test
    # shapiro0 = stats.shapiro(group0)
    # shapiro1 = stats.shapiro(group1)
```

```

# Here we just do both t-test and Mann-Whitney and report them.
t_stat, t_p = stats.ttest_ind(group0, group1, equal_var=False)
u_stat, u_p = stats.mannwhitneyu(group0, group1, alternative="two-sided")

results.append({
    "feature": col,
    "t_statistic": t_stat,
    "t_pvalue": t_p,
    "mannwhitney_u": u_stat,
    "mannwhitney_pvalue": u_p,
    "mean_class0": group0.mean(),
    "mean_class1": group1.mean()
})

stats_df = pd.DataFrame(results).set_index("feature")
stats_df.sort_values("mannwhitney_pvalue", inplace=True)
stats_df

```

Out[22]:

		t_statistic	t_pvalue	mannwhitney_u	mannwhitney_pvalue	mean_clas
	feature					
	Insulin	1.064205	0.287303	2041508.0	0.074129	8.7746
	Resistin	-1.665580	0.095878	1915307.5	0.090990	12.8357
	Age	0.504255	0.614111	1994986.5	0.613757	56.3688
	Leptin	0.311227	0.755645	1993972.0	0.633722	25.2204
	BMI	0.590046	0.555195	1992751.0	0.657865	27.4682
	MCP.1	-0.379380	0.704427	1961162.5	0.669250	510.4892
	Glucose	0.309659	0.756837	1985751.5	0.800449	114.0173
	HOMA	0.784348	0.432885	1985021.0	0.818131	2.0468
	Adiponectin	0.471932	0.637003	1969372.5	0.840666	9.4039



In []:

```

corr = df[feature_cols + [TARGET_COL]].corr(method="pearson")
corr

```

Out[]:

	Age	BMI	Glucose	Insulin	HOMA	Leptin	Adiponectin
Age	1.000000	0.019139	0.002611	0.000742	0.024631	0.005058	-0.000100
BMI	0.019139	1.000000	-0.014469	0.011574	0.012795	-0.011482	-0.000100
Glucose	0.002611	-0.014469	1.000000	0.000558	-0.018232	-0.006461	-0.000100
Insulin	0.000742	0.011574	0.000558	1.000000	-0.013142	0.005678	-0.000100
HOMA	0.024631	0.012795	-0.018232	-0.013142	1.000000	0.006751	-0.000100
Leptin	0.005058	-0.011482	-0.006461	0.005678	0.006751	1.000000	-0.000100
Adiponectin	-0.013493	-0.002348	0.016363	-0.014560	0.003433	0.024054	-0.000100
Resistin	-0.013438	0.005303	0.014179	0.036136	0.000272	0.014217	-0.000100
MCP.1	-0.019410	-0.023389	0.001837	-0.011382	0.009965	-0.003359	-0.000100
Classification	-0.007965	-0.009352	-0.004893	-0.016826	-0.012441	-0.004914	-0.000100

6. Correlation Analysis

A Pearson correlation matrix and heatmap were generated.

Key Findings

- All correlations are extremely low ($|r| < 0.03$ in most cases).
- No meaningful linear relationships exist between features.
- Correlation between each feature and `Classification` is close to zero.

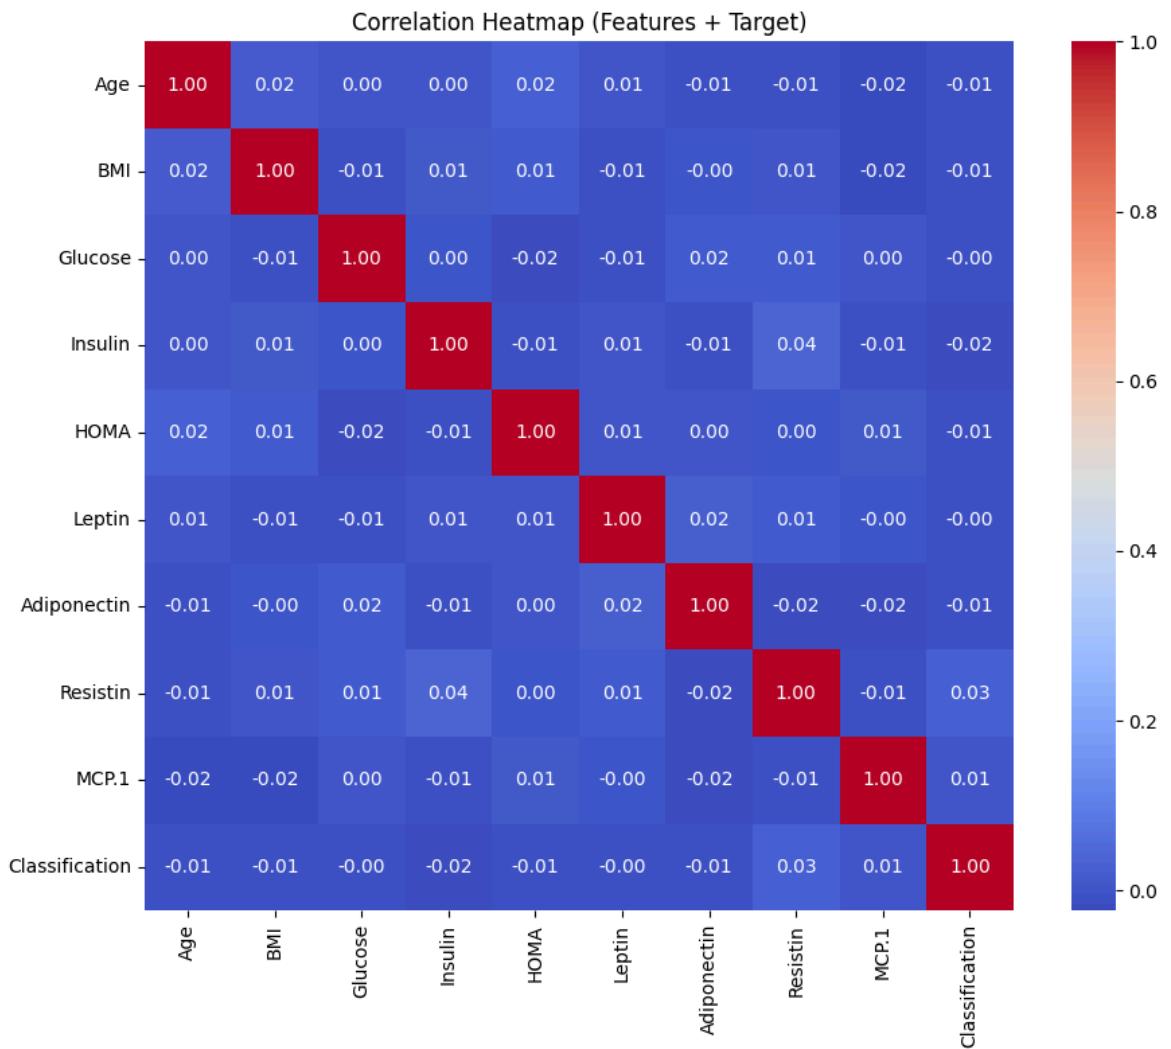
Interpretation

The dataset lacks linear structure.

Effective models must capture **nonlinear** relationships or feature interactions.

In [24]:

```
plt.figure(figsize=(10, 8))
sns.heatmap(corr, annot=True, fmt=".2f", cmap="coolwarm", square=True)
plt.title("Correlation Heatmap (Features + Target)")
plt.tight_layout()
plt.show()
```



7. Outlier Detection

Z-score analysis flagged outliers using the threshold $|z| > 3$:

Feature	Outlier Count
Insulin	110
HOMA	112
Resistin	43
Others	0

Interpretation

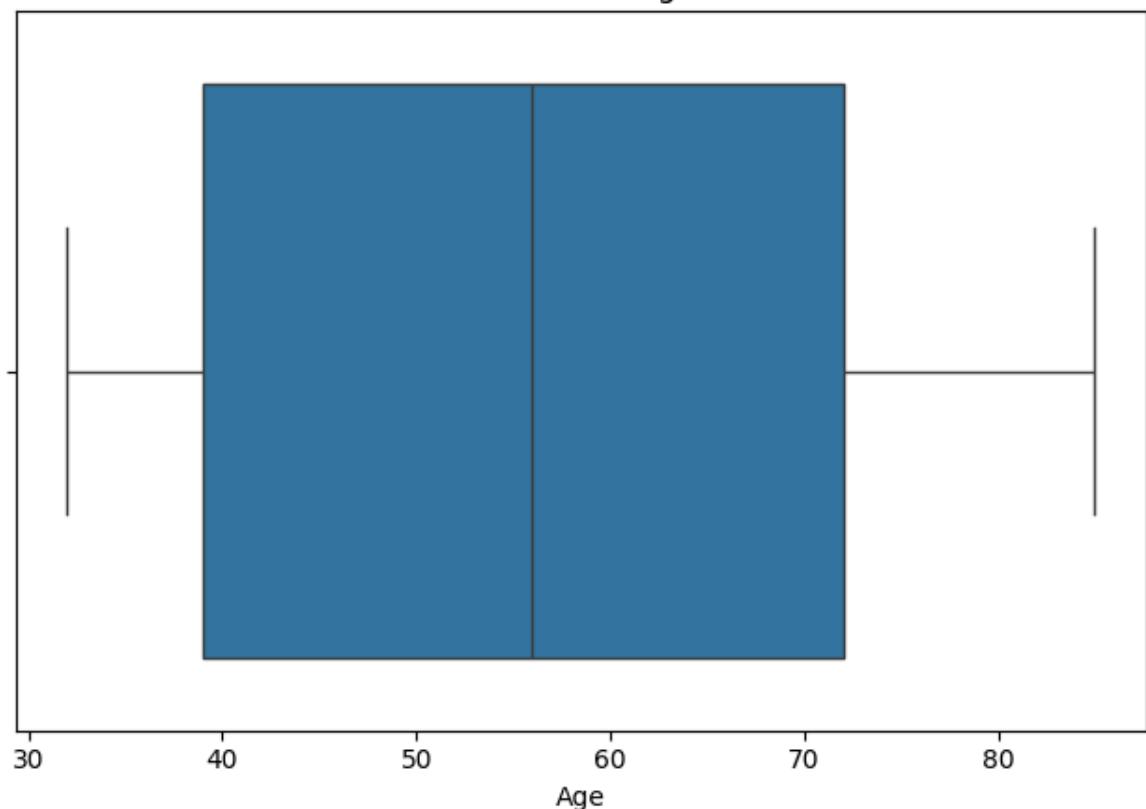
Outliers are concentrated in metabolic markers, which commonly display high biological variability.

Given the dataset size, retaining them is reasonable, but scaling becomes crucial.

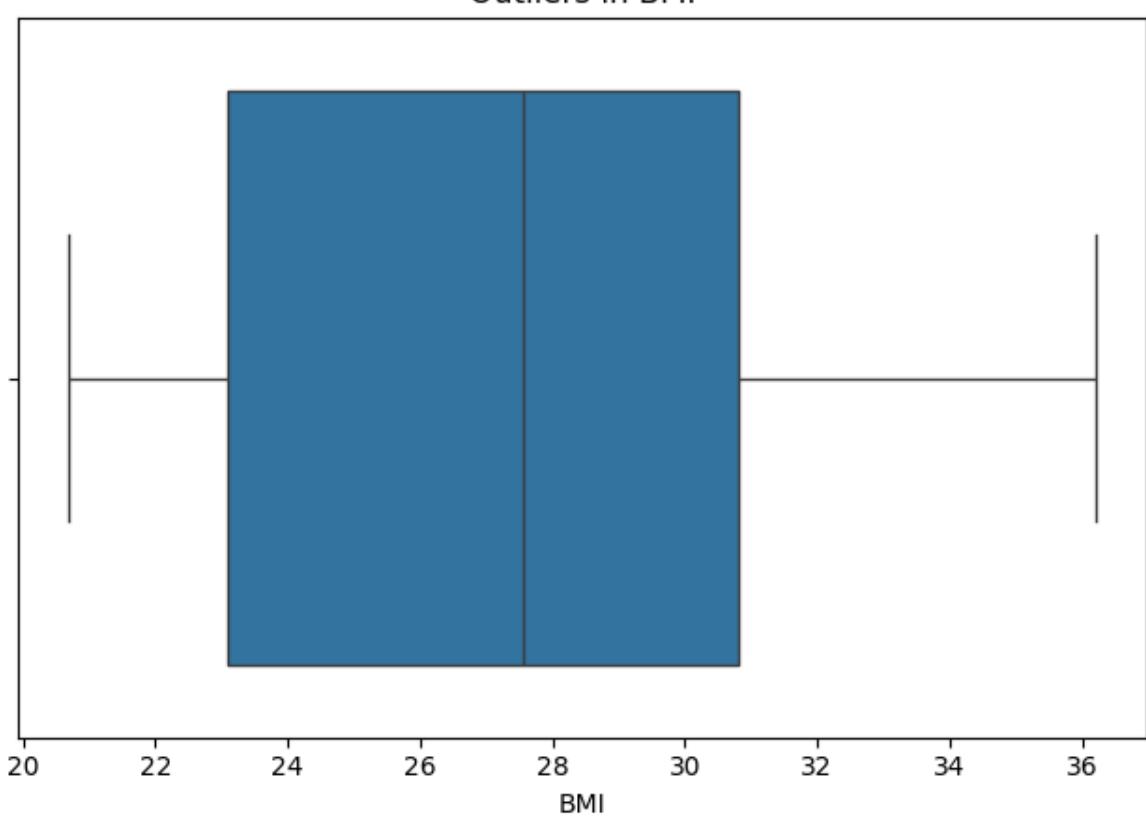
```
In [ ]: for col in feature_cols:
    fig, ax = plt.subplots()
    sns.boxplot(x=df[col], ax=ax)
    ax.set_title(f"Outliers in {col}")
```

```
plt.tight_layout()  
plt.show()
```

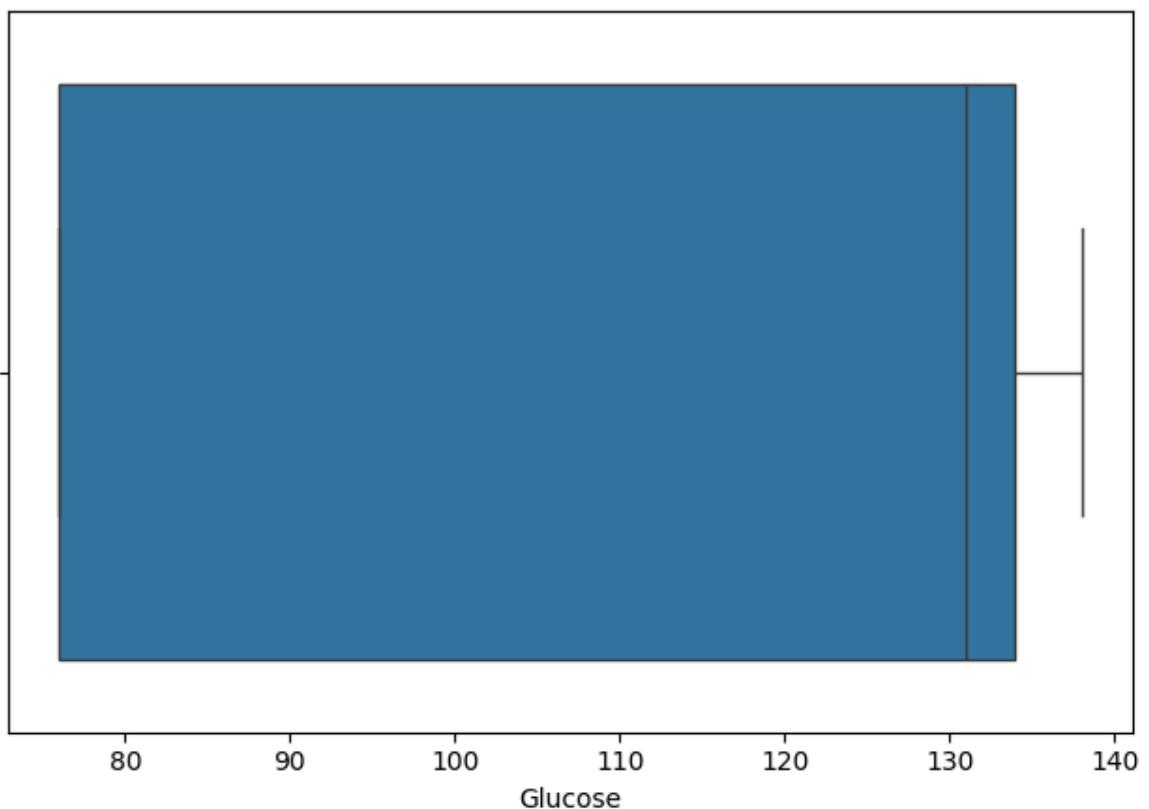
Outliers in Age



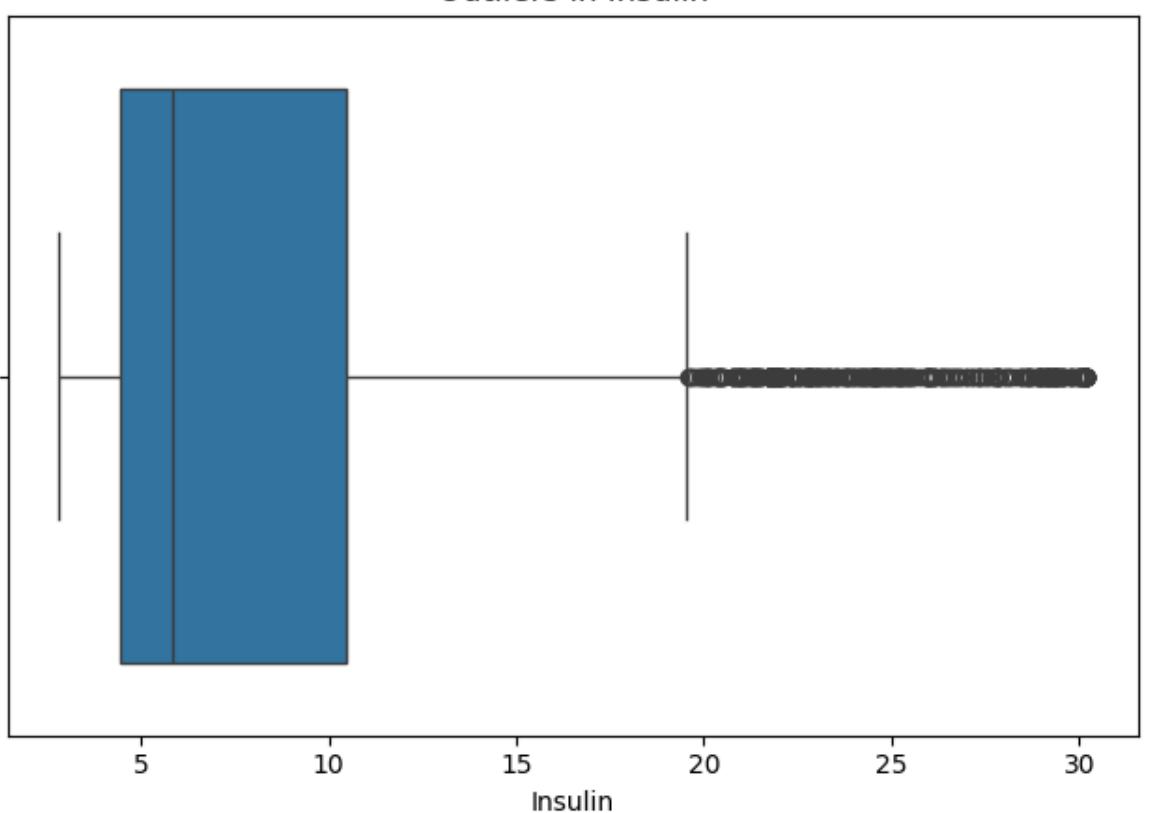
Outliers in BMI



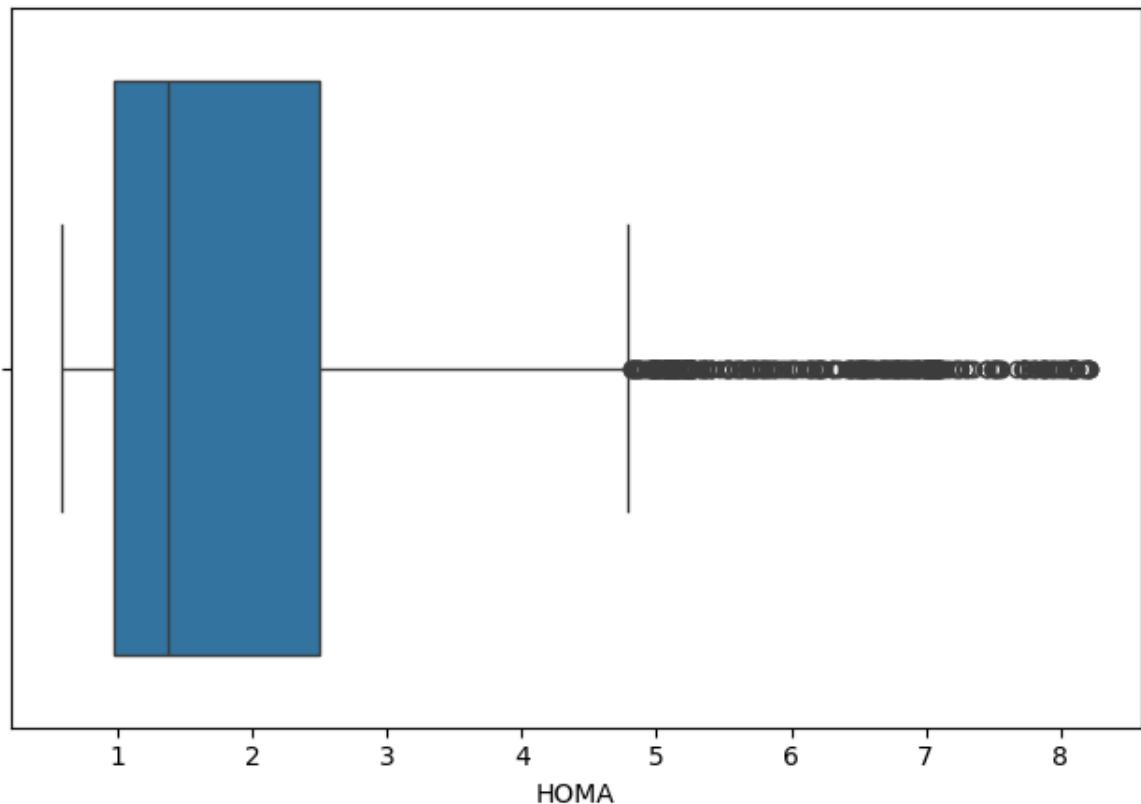
Outliers in Glucose



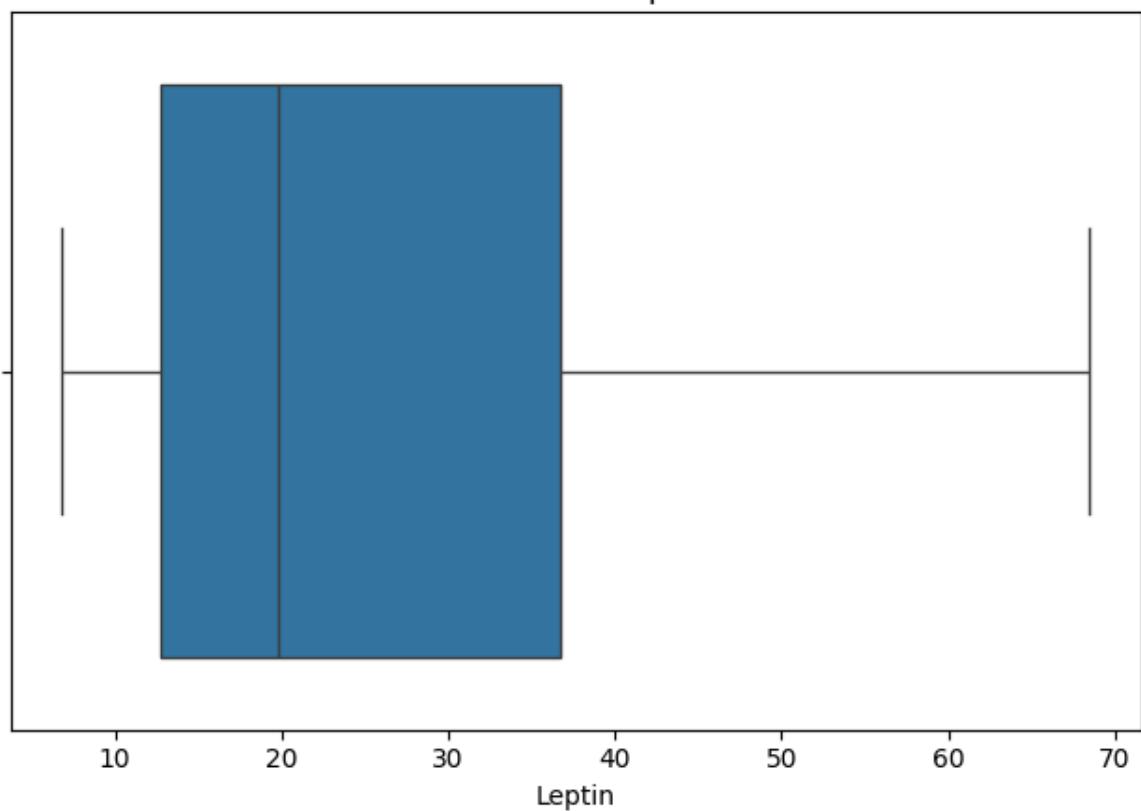
Outliers in Insulin



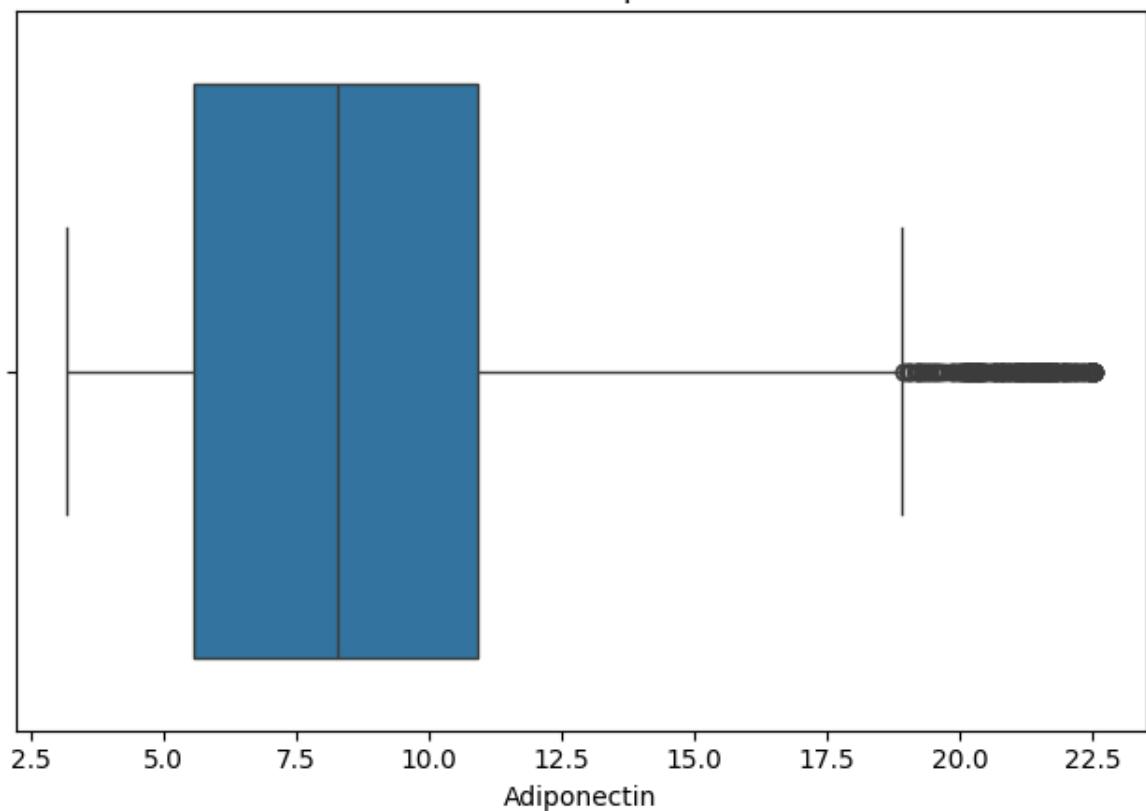
Outliers in HOMA



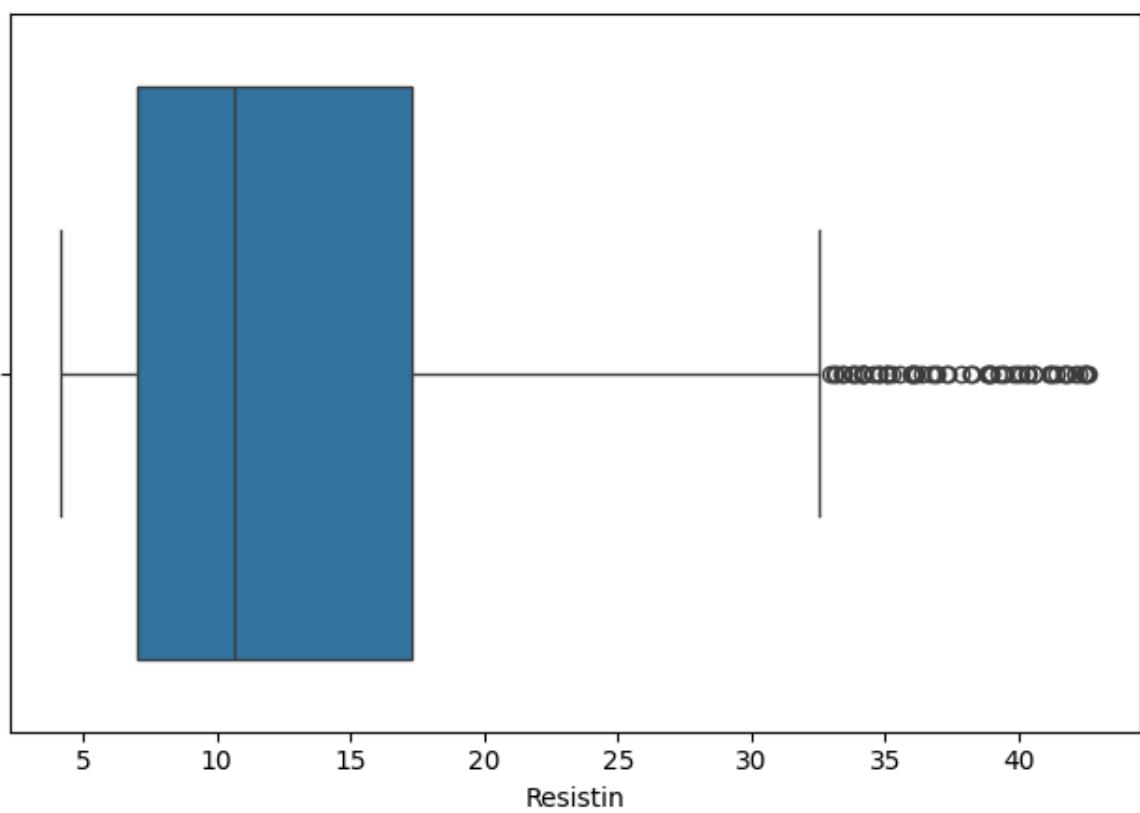
Outliers in Leptin



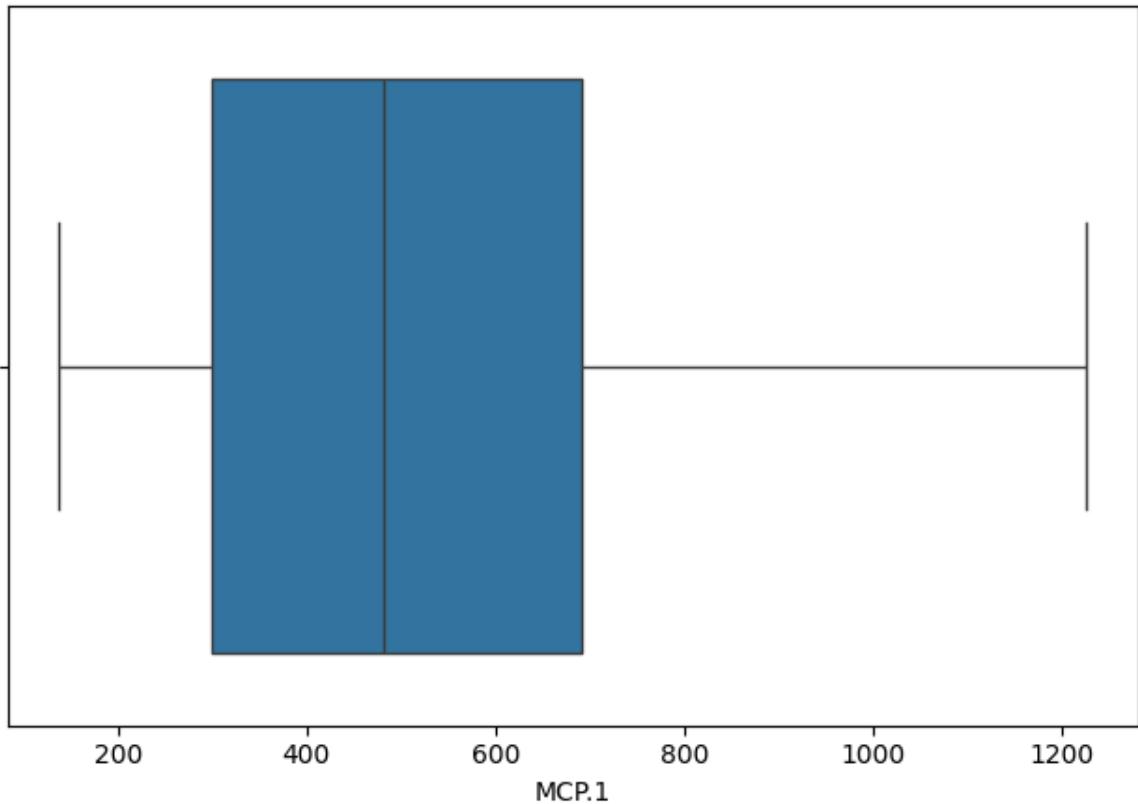
Outliers in Adiponectin



Outliers in Resistin



Outliers in MCP.1



```
In [26]: z_scores = np.abs(stats.zscore(df[feature_cols]))
z_df = pd.DataFrame(z_scores, columns=feature_cols)

# Consider |z| > 3 as outlier
outlier_mask = z_df > 3
outlier_counts = outlier_mask.sum()

print("Outlier counts per feature (|z| > 3):")
outlier_counts
```

Outlier counts per feature ($|z| > 3$):

```
Out[26]: Age          0
          BMI         0
          Glucose      0
          Insulin     110
          HOMA        112
          Leptin       0
          Adiponectin  0
          Resistin     43
          MCP.1        0
          dtype: int64
```

8. Scaling Effects

StandardScaler was applied to all numerical features.

Results

- All features successfully standardized to ~ 0 mean and unit variance.
- Skewed features remain skewed (scaling does not normalize distribution).

Why scaling matters

Essential for algorithms such as:

- Logistic Regression
- SVM
- kNN
- Neural Networks

Particularly important because the Coimbra dataset contains **wide variance differences** across metabolic markers.

In [27]:

```
scaler = StandardScaler()
scaled_features = scaler.fit_transform(df[feature_cols])
scaled_df = pd.DataFrame(scaled_features, columns=feature_cols)

print("Original feature summary:")
display(df[feature_cols].describe().T)

print("\nScaled feature summary (mean ~0, std ~1):")
display(scaled_df.describe().T)
```

Original feature summary:

	count	mean	std	min	25%	50%
Age	4000.0	56.210750	17.809650	32.000000	39.000000	56.000000
BMI	4000.0	27.422280	4.413884	20.690751	23.079053	27.558485
Glucose	4000.0	113.876500	25.837795	76.000000	76.000000	131.000000
Insulin	4000.0	8.654001	6.435160	2.821000	4.421750	5.818000
HOMA	4000.0	2.024332	1.625638	0.590033	0.970090	1.373842
Leptin	4000.0	25.137737	15.096446	6.831900	12.712750	19.805050
Adiponectin	4000.0	9.364896	4.674244	3.192272	5.580210	8.286938
Resistin	4000.0	13.053667	7.454424	4.190320	7.022095	10.692780
MCP.1	4000.0	512.183456	253.279015	137.488000	299.665000	482.308000

Scaled feature summary (mean ~0, std ~1):

		count	mean	std	min	25%	50%	75%
Age	4000.0	1.523226e-16	1.000125	-1.359587	-0.966493	-0.011835	0.886667	0.886667
BMI	4000.0	-3.144152e-16	1.000125	-1.525271	-0.984115	0.030862	0.768724	0.768724
Glucose	4000.0	2.664535e-16	1.000125	-1.466117	-1.466117	0.662814	0.778937	0.778937
Insulin	4000.0	-1.509903e-16	1.000125	-0.906540	-0.657758	-0.440759	0.281652	0.281652
HOMA	4000.0	6.572520e-17	1.000125	-0.882410	-0.648591	-0.400194	0.294348	0.294348
Leptin	4000.0	3.197442e-17	1.000125	-1.212744	-0.823143	-0.353285	0.764018	0.764018
Adiponectin	4000.0	1.616485e-16	1.000125	-1.320726	-0.809791	-0.230645	0.332373	0.332373
Resistin	4000.0	-6.394885e-17	1.000125	-1.189154	-0.809228	-0.316749	0.569006	0.569006
MCP.1	4000.0	3.126388e-16	1.000125	-1.479563	-0.839173	-0.117969	0.707328	0.707328



9. Feature Importance (Mutual Information & ANOVA F-test)

Mutual Information Rankings

1. **HOMA** (0.017)
2. **Resistin** (0.006)
3. **Glucose** (0.004)
4. **Adiponectin** (0.003)
5. Others: 0.0

ANOVA F-test Rankings

- **Resistin**: $F = 2.75$, $p \approx 0.097$
- Insulin: $F \approx 1.13$
- Others: $F < 1$

Interpretation

- No feature demonstrates strong predictive value individually.
- HOMA and Resistin contain the most information, but signals are weak.
- Models must rely on **feature combinations** and **nonlinear effects**.

```
In [28]: X = df[feature_cols].values
y = df[TARGET_COL].values

mi = mutual_info_classif(X, y, random_state=42)

mi_df = pd.DataFrame({
    "feature": feature_cols,
    "mutual_information": mi
}).sort_values("mutual_information", ascending=False).set_index("feature")

mi_df
```

Out[28]: **mutual_information**

feature	
HOMA	0.017155
Resistin	0.006323
Glucose	0.004276
Adiponectin	0.003197
Age	0.000000
BMI	0.000000
Insulin	0.000000
Leptin	0.000000
MCP1	0.000000

```
In [29]: f_vals, f_pvals = f_classif(X, y)

f_df = pd.DataFrame({
    "feature": feature_cols,
    "f_value": f_vals,
    "p_value": f_pvals
}).sort_values("f_value", ascending=False).set_index("feature")

f_df
```

Out[29]:

feature	f_value	p_value
Resistin	2.753596	0.097114
Insulin	1.132220	0.287367
HOMA	0.618858	0.431519
BMI	0.349664	0.554337
Age	0.253670	0.614530
Adiponectin	0.224822	0.635416
MCP.1	0.144050	0.704308
Leptin	0.096552	0.756023
Glucose	0.095709	0.757057

Final Summary

Key Findings

1. Feature distributions show strong skewness in Insulin, HOMA, Resistin.
2. Healthy vs Patient groups have heavily overlapping distributions.
3. Statistical tests show weak class separation.
4. Correlations are extremely low across all feature pairs.
5. Outliers exist mainly in metabolic markers.
6. Scaling is required due to differing magnitude ranges.
7. Feature importance scores are very small, indicating weak individual predictors.

Overall Conclusion

The Coimbra dataset does **not** contain strong linear signals in single features. Successful classification will require:

- Nonlinear models (Random Forest, Gradient Boosting, SVM RBF)
- Use of interaction terms
- Combined feature effects
- Robust preprocessing and scaling

This dataset is best modeled using **multivariate, nonlinear techniques** rather than relying on any one biomarker alone.